



VPRO SKILLS

FAQ / INTERVIEW PREPARATION

Agents Interview Questions & Answers

STUDENT LEARNING RESOURCE

Instagram: [vpro.skill](#). | YouTube: [@vpro.skills](#)

Contact: 9010001847

Contents

Agents 4

Quick Revision 131

How to Use This Guide

This handbook brings together 218 curated questions and answers on Agents, organized into a single quick-reference list for focused revision. Each question is numbered for easy reference, and answers are broken down into definitions, examples, code, and comparison tables wherever the source material provides them.

- Code blocks are shown in a monospaced, shaded panel — read them alongside the explanation above.
- Comparison tables summarize key differences at a glance; use them for quick revision before an interview.
- The Quick Revision page at the end lists every question so you can self-test your recall.

AGENTS

Q001. What is an AI Agent?

SIMPLE ANSWER

An AI Agent is an AI system that can understand a goal, decide what to do, use tools, and take actions to complete the task.

A normal LLM mainly generates an answer.

An Agent can think about the task, use tools, take actions, check the result, and continue until the task is completed.

EXAMPLE

If I ask:

"Find the cheapest flight and book it."

An Agent can:

- Search flights.
- Compare prices.
- Select an option based on the given rules.
- Ask for confirmation if required.
- Make the booking.

Q002. What is the difference between an LLM and an AI Agent?

SIMPLE ANSWER

An LLM generates text or answers based on the prompt. An AI Agent uses an LLM as its brain and can also use tools, memory, and actions to complete a task.

EXAMPLE

LLM:

User: "What is today's weather?"

LLM → Gives an answer based on available information.

Agent:

User: "Check today's weather and tell me whether I should carry an umbrella."

Agent → Calls a weather tool → gets current weather → analyzes it → gives a recommendation.

INTERVIEW LINE

"I consider the LLM as the reasoning engine, while the Agent is the complete system around the LLM that can reason, use tools and perform actions."

Q003. What are the main components of an AI Agent?

An Agent normally contains:

- LLM – Brain/reasoning
- Prompt – Instructions
- Tools – External capabilities
- Memory – Stores useful information
- State – Current execution information
- Planning/Reasoning – Decides what to do
- Orchestration – Controls the workflow

Simple architecture:

```
User
  ↓
AI Agent
  ↓
LLM
  ↓
Decision
  ↓
Tools
  ↓
External Systems
  ↓
Result
  ↓
LLM
  ↓
Final Answer
```

Q004. How does an AI Agent make decisions?

SIMPLE ANSWER

The Agent sends the user's request and available information to the LLM. The LLM decides what action is required. If a tool is needed, the Agent calls that tool and uses the result for the next step.

EXAMPLE

User:

"What is the price of Bitcoin?"

Agent:

```
User Question
  ↓
LLM
  ↓
Need current price?
  ↓
Call Market API
  ↓
Get Price
  ↓
LLM
  ↓
Final Answer
```

INTERVIEW LINE

"The LLM decides the next action, while the Agent framework executes that action."

Q005. What is the Agent Execution Loop?

This is a very important interview question.

An Agent usually works in a loop:

```
Understand (LLM)
  ↓
Plan (tool selection)
  ↓
Act (execute)
  ↓
Observe (result verification)
  ↓
Decide Next Step
  ↓
Act Again
  ↓
...
  ↓
Final Answer
```

EXAMPLE

Suppose we ask:

"Find the weather in Hyderabad and suggest what I should wear."

The Agent may do:

```
Understand request
  ↓
Call weather tool
  ↓
Get weather
  ↓
Analyze result
  ↓
Generate recommendation
  ↓
Final Answer
```

If multiple actions are required, the Agent continues the loop.

Experienced answer:

"An Agent typically follows an execution loop where it receives a goal, decides the next action, executes a tool or operation, observes the result, updates its state, and continues until the task is completed or a stopping condition is reached."

Q006. What is Tool Calling?

SIMPLE ANSWER

Tool calling means allowing an LLM or Agent to use an external function or service to perform a task.

The tool could be:

- Weather API
- Database
- Search engine
- Calculator
- Email service
- CRM
- Python function
- Internal company API

EXAMPLE

```
def get_weather(city):
    # call weather API
    return weather
```

The Agent can decide:

"I need current weather, so I should call get_weather()."

Q007. What is Function Calling?

SIMPLE ANSWER

Function calling is a mechanism where the LLM returns the name of a function and the required arguments so the application can execute that function.

For example:

User:

"Get weather for Hyderabad."

LLM may produce something conceptually like:

```
Function: get_weather
Arguments:
  city = "Hyderabad"
```

Your application executes the function and sends the result back to the LLM.

Important:

The LLM usually does not directly execute your Python function.

Your application/framework executes the function.

INTERVIEW LINE

"Function calling provides the structured interface between the LLM and application tools."

Q008. What is the difference between Tool Calling and Function Calling?

In simple terms, they are closely related.

Function calling generally refers to the structured mechanism where a model requests a function with arguments.

Tool calling is the broader concept of allowing an Agent to use external capabilities such as functions, APIs, databases or services.

In modern Agent frameworks, you will often see the terms used almost interchangeably.

Q009. What is Memory in an AI Agent?

SIMPLE ANSWER

Memory allows an Agent to remember information that may be useful later.

EXAMPLE

User says:

"My name is Ravi."

Later:

"What is my name?"

If the Agent has appropriate memory, it can answer:

"Your name is Ravi."

Memory is useful for:

- Conversations
- User preferences
- Previous tasks
- Important information
- Long-running workflows

Q010. What is Short-Term Memory?

SIMPLE ANSWER

Short-term memory contains information needed during the current conversation or task.

EXAMPLE

```
User: My name is Ravi.  
User: I work with Python.  
User: What technology do I use?
```

The Agent can use the current conversation context to answer:

"You work with Python."

This information may not necessarily be stored permanently.

Q011. What is Long-Term Memory?

SIMPLE ANSWER

Long-term memory stores information that can be reused across different conversations or sessions.

EXAMPLE

```
User preference:  
Preferred language = Python
```

Next week, the Agent may retrieve that information when needed.

Long-term memory could be implemented using:

- Databases
- Vector databases
- Key-value stores
- Other persistent storage systems

Q012. What is Planning in Agentic AI?

SIMPLE ANSWER

Planning means breaking a large task into smaller steps and deciding the order in which those steps should be performed.

EXAMPLE

User:

"Build and deploy a web application."

The Agent may create a plan:

1. Understand requirements
2. Create backend
3. Create frontend
4. Test application
5. Build Docker image
6. Deploy
7. Verify deployment

The Agent can then execute these steps.

Q013. What is Reasoning in an AI Agent?

SIMPLE ANSWER

Reasoning means using the available information to decide what action should be taken next.

EXAMPLE

User:

"Find employees who joined after 2025."

The Agent understands that it needs employee data.

It may decide:

```
Need employee information
  ↓
Use database tool
  ↓
Run query
  ↓
Analyze result
  ↓
Return employees
```

INTERVIEW LINE

"Reasoning helps the Agent decide the next action based on the current goal and available information."

Q014. What is Reflection?

SIMPLE ANSWER

Reflection means the Agent checks its own result and tries to identify mistakes or improvements.

EXAMPLE

Agent generates:

Report generated.

Then it checks:

Is the report complete?
Are important sections missing?
Are there errors?

If something is wrong, it can improve the result.

Simple flow:

```
Generate
  ↓
Check
  ↓
Find Problems
  ↓
Improve
  ↓
Final Result
```

Q015. What is ReAct?

ReAct means:

Reason + Act

The basic idea is that the Agent alternates between deciding what to do and taking actions.

Conceptually:

```
Reason
↓
Action
↓
Observation
↓
Reason
↓
Action
↓
Observation
↓
Final Answer
```

EXAMPLE

User:

"What is the population of the capital of India?"

Agent:

```
Reason → Need to identify capital
Action → Search/knowledge tool
Observation → New Delhi
Reason → Need population
Action → Search population
Observation → Result
Final Answer
```

INTERVIEW LINE

"ReAct combines reasoning and tool-based actions so the Agent can use information obtained from one step to decide the next step."

Q016. What is the difference between a Workflow and an Agent?

This is very important for experienced interviews.

WORKFLOW

A workflow usually follows predefined steps.

```
Step 1
↓
Step 2
↓
Step 3
↓
Step 4
```

The developer controls the flow.

Agent

An Agent can decide what to do next based on the task.

```
User
↓
Agent
↓
Decision
```

↓ □

```
Tool Tool Tool
```

Simple example

WORKFLOW

```
Read PDF
↓
Extract Text
↓
Create Summary
```

Agent:

```
Read PDF
↓
Agent decides:
"Do I need search?"
↓
Search
↓
Agent decides:
"Do I need another tool?"
↓
...
```

Experienced answer:

"A workflow is generally deterministic because the developer defines the execution path. An Agent is more dynamic because the model can decide the next action. In production, I often combine both: deterministic workflows for critical business logic and Agents where dynamic decision-making is useful."

Q017. When should we use an Agent instead of a normal LLM application?

Use an Agent when the application needs:

- Multiple steps
- Dynamic decision-making
- Tool usage
- External APIs
- Database access

- Planning
- Conditional actions
- Iterative problem solving

EXAMPLE

A chatbot:

Question → LLM → Answer

You may not need an Agent.

But:

```
graph TD
    A[User Request] --> B[Search Database]
    B --> C[Check API]
    C --> D[Analyze Data]
    D --> E[Send Email]
    E --> F[Generate Report]
```

An Agent architecture may be useful.

Experienced answer:

"I don't use Agents for every LLM application. If the workflow is simple and predictable, I prefer a normal LLM pipeline. I use Agents when the system needs dynamic tool selection or multi-step decision-making."

Q018. What are the limitations of AI Agents?

Important limitations include:

1. Cost — Agents may call the LLM many times.
2. Latency — Multiple tool calls and LLM calls can make the application slower.
3. Hallucination — The Agent may make an incorrect decision.
4. Unpredictability — The Agent may choose an unexpected path.
5. Security — Agents can access tools and systems, so permissions must be controlled.
6. Debugging — Complex Agent workflows can be difficult to debug.

INTERVIEW ANSWER

"The main challenges are cost, latency, reliability, security and unpredictable behavior. For production systems, I control the Agent using tool permissions, validation, timeouts, retries, state management, logging and human approval for sensitive actions."

Q019. How do you prevent an Agent from entering an infinite loop?

SIMPLE ANSWER

I define clear stopping conditions and limits.

For example:

```
Maximum iterations = 10  
Maximum execution time = 60 seconds
```

Other techniques:

- Maximum tool calls
- Timeout
- Retry limits
- State validation
- Error handling
- Explicit END condition
- Human approval for sensitive workflows

Experienced answer:

"I never allow an Agent to run without limits in production. I normally configure maximum iterations, timeouts, retry limits and explicit termination conditions."

Q020. How do you handle hallucinations in Agents?

SIMPLE ANSWER

I reduce hallucination by giving the Agent reliable information and controlling how it can take actions.

Common techniques:

- RAG
- Trusted tools
- Database queries
- Structured outputs
- Validation
- Grounding
- Human approval
- Guardrails
- Evaluation

EXAMPLE

Instead of asking:

"What is the customer's balance?"

allow the Agent to call:

```
get_customer_balance()
```

Then the Agent uses the actual database result.

Experienced answer:

"I don't depend only on the LLM's knowledge for factual business data. I connect the Agent to trusted sources such as databases, APIs or RAG systems and validate important outputs before taking actions."

Q021. How do you evaluate an AI Agent?

SIMPLE ANSWER

I evaluate both the final answer and the Agent's behavior.

I check:

Answer quality

- Is the answer correct?
- Is it relevant?
- Is it grounded?

Agent behavior

- Did it select the correct tool?
- Did it use the correct arguments?
- Did it follow the expected workflow?
- Did it stop correctly?
- How many LLM calls did it make?
- How many tool calls did it make?
- How much time did it take?
- How much did it cost?

Experienced answer:

"For Agent evaluation, I look beyond the final response. I evaluate tool selection, tool arguments, execution path, task completion, latency, token usage, cost and final answer quality."

Q022. What is an Agent Loop?

This question is often asked separately.

A simple Agent Loop is:

```
User Goal
  ↓
LLM
  ↓
Decide
  ↓
Tool
  ↓
Result
  ↓
LLM
  ↓
Decide Again
  ↓
Tool
  ↓
Result
  ↓
Final Answer
```

The important idea is:

The Agent doesn't necessarily perform only one action. It can observe the result and decide what to do next.

Q023. What is State in an Agent?

SIMPLE ANSWER

State is the information the Agent needs to know about the current task.

For example:

```
state = {
    "user_query": "...",
    "messages": [],
    "search_results": [],
    "current_step": "...",
    "final_answer": None
}
```

As the Agent works, the state changes.

EXAMPLE

```

Initial State
  ↓
Search Result Added
  ↓
Analysis Added
  ↓
Validation Added
  ↓
Final Answer Added


```

This becomes especially important when working with LangGraph.

Q024. What is Agent Orchestration?**SIMPLE ANSWER**

Orchestration means controlling how different Agents, tools and tasks work together.

For example:



The orchestration layer decides:

"Which Agent should work next?"


Q025. How do you design a production-ready AI Agent?

This is a very strong experienced-level question.

You can answer:

"First, I clearly define the Agent's responsibility and available tools. Then I design the state and workflow. I add validation, error handling, retries, timeouts and access control. I keep secrets outside the code. I add logging and monitoring to track LLM calls and tool calls. Finally, I evaluate the Agent using real test cases and deploy it using containers and CI/CD."

A good production architecture is:



Question	One-line answer
What is Agent?	AI system that can reason, use tools and take actions.
LLM vs Agent?	LLM generates; Agent can decide and act.
Tool Calling?	Agent uses an external capability.
Function Calling?	LLM requests a function with structured arguments.
Memory?	Stores information useful for future steps/conversations.
Planning?	Breaking a complex goal into smaller steps.
Reasoning?	Deciding what action should happen next.
Reflection?	Checking and improving the result.
ReAct?	Reason + Act + Observe loop.
Workflow?	Developer-defined execution path.
Agent?	Dynamic decision-making system.
State?	Information about the current task.
Agent Loop?	Decide → Act → Observe → Decide again.
Tool?	External capability available to the Agent.
Hallucination?	Agent/LLM produces unsupported or incorrect information.

Question	One-line answer
Guardrails?	Rules that restrict and validate Agent behavior.
Agent Evaluation?	Measure answer quality and Agent behavior.
Orchestration?	Controlling how Agents, tools and tasks work together.

Q026. LangChain – Interview Questions

Q027. What is LangChain?

SIMPLE ANSWER

LangChain is a framework used to build applications with Large Language Models (LLMs).

It provides ready-made components for:

- LLMs
- Prompts
- Tools
- Agents
- RAG
- Memory
- Document loaders
- Retrievers
- Structured output

Instead of writing everything from scratch, we can use LangChain components.

INTERVIEW ANSWER

"LangChain is a framework that helps us build LLM-based applications by providing integrations for models, prompts, tools, retrievers, agents and other components."

Q028. Why do we need LangChain?

Suppose we want to build an AI application.

Without LangChain:


```

Our Code
  ↓
OpenAI API
  ↓
Database
  ↓
PDF Processing
  ↓
Tools
  ↓
Custom Logic

```

We need to write and maintain many integrations ourselves.

With LangChain:



INTERVIEW ANSWER

"LangChain reduces the amount of custom code required to connect LLMs with tools, data sources and application logic."

Q029. What are the main components of LangChain?

The important components are:

```

LLM
  ↓
Prompt
  ↓
Chain / Runnable
  ↓
Tool
  ↓
Agent
  ↓
Retriever
  ↓
Memory / State

```

Some common components are:

Models

Connect to LLMs.

EXAMPLES

- OpenAI
- Anthropic
- Google
- Azure OpenAI

Prompts

Define instructions for the LLM.

Tools

Allow the Agent to perform actions.

Retrievers

Find relevant information from a knowledge source.

Document Loaders

Load:

- PDF
- Word
- CSV
- Web pages
- Text files

Agents

Allow the LLM to decide which tools to use.

Q030. What is a Prompt Template?

A Prompt Template is a reusable prompt with variables.

Instead of writing:

```
Explain Python in simple English.
```

every time, we can create:

```
Explain {topic} in simple English.
```

Then:

```
topic = "LangChain"
```

The final prompt becomes:

```
Explain LangChain in simple English.
```

INTERVIEW ANSWER

"A Prompt Template allows us to create reusable prompts with dynamic variables."

Q031. What is an LLM in LangChain?

The LLM is the brain of the application.

For example:

```
User
  ↓
LangChain
  ↓
OpenAI / Claude / Gemini
  ↓
Response
```

LangChain provides a common way to interact with different models.

INTERVIEW ANSWER

"LangChain provides abstractions that allow applications to work with different LLM providers without tightly coupling the application logic to one provider."

Q032. What is a Chain?

A Chain means connecting multiple operations together.

For example:

```
User Question
  ↓
Prompt
  ↓
LLM
  ↓
Output
```

Another example:

```
Question
  ↓
Retriever
  ↓
Documents
  ↓
Prompt
  ↓
LLM
  ↓
Answer
```

This is commonly used in RAG applications.

Simple definition

A Chain connects multiple steps to complete a task.

Q033. What is LCEL?

LCEL means:

LangChain Expression Language

It provides a way to connect LangChain components together.

For example:

```
Prompt → LLM → Parser
```

Conceptually:

```
chain = prompt | llm | parser
```

The | means:

Send the output of one component to the next component.

INTERVIEW ANSWER

"LCEL is LangChain's expression language for composing components into runnable pipelines."

Q034. What is Runnable?

A Runnable is a component that can be executed as part of a LangChain pipeline.

For example:

```
Prompt  
↓  
LLM  
↓  
Parser
```

Each component can participate in the pipeline.

Simple interview answer

"Runnable is a common interface used by LangChain components so they can be composed and executed together."

Q035. What is a Tool in LangChain?

A Tool is an external capability that an Agent can use.

EXAMPLES

Calculator
Weather API
Database
Search
Python Function
Email API
CRM API

For example:

```
def get_customer(customer_id):  
    # database operation  
    return customer
```

We can expose this function as a tool.

The Agent can then decide:

"I need customer information, so I should use the customer tool."

Q036. What is an Agent in LangChain?

A LangChain Agent uses an LLM to decide what action to take.

EXAMPLE

```
User  
↓  
Agent  
↓  
LLM  
↓  
Need weather?  
↓  
Weather Tool  
↓  
Result  
↓  
LLM  
↓  
Final Answer
```

INTERVIEW ANSWER

"A LangChain Agent uses an LLM to decide which tools to use and in what sequence to complete a task."

Q037. What is the difference between Chain and Agent?

This is a very common interview question.

Chain

The developer defines the flow.

```
Step 1
↓
Step 2
↓
Step 3
```

Agent

The LLM decides the next action.

```
Agent
↓
LLM Decision
/ | \
Tool A Tool B Tool C
```

INTERVIEW ANSWER

"A Chain follows a predefined sequence, while an Agent can dynamically decide which tool or action to take."

Q038. How does a LangChain Agent work?

Suppose we ask:

"What is the weather in Hyderabad?"

The process can be:

```
User Question
↓
Agent
↓
LLM
↓
Select Weather Tool
↓
Weather API
↓
Weather Result
↓
LLM
↓
Final Answer
```

The important point:

The LLM decides which tool is required.

The LangChain runtime/application executes the tool.

Q039. What is a Retriever?

A Retriever is responsible for finding relevant information.

For example:

```
User Question
      ↓
Retriever
      ↓
Relevant Documents
```

In RAG:

```
Question
      ↓
Embedding
      ↓
Retriever
      ↓
Relevant Chunks
      ↓
LLM
      ↓
Answer
```

INTERVIEW ANSWER

"A retriever searches a knowledge source and returns the information relevant to the user's query."

Q040. What is RAG in LangChain?

RAG means:

Retrieval-Augmented Generation

Instead of asking the LLM to answer only from its training knowledge, we provide relevant information from our own data.

EXAMPLE

Suppose a company has:

```
Company Policies.pdf
Employee Handbook.pdf
Leave Policy.pdf
```

User asks:

"How many days of annual leave do employees get?"

RAG:

```
graph TD;
    Question --> Retriever;
    Retriever --> Relevant[Relevant PDF chunks];
    Relevant --> Prompt[Prompt + Documents];
    Prompt --> LLM;
    LLM --> Answer;
```

INTERVIEW ANSWER

"RAG allows an LLM to answer questions using external or private data retrieved at query time."

Q041. What is a Document Loader?

A Document Loader loads data into LangChain.

EXAMPLES

```
graph TD;
    PDF --> Loader[Document Loader];
    Loader --> Documents;
```

Common sources:

- PDF
- TXT
- CSV
- Word
- Web pages
- Databases

INTERVIEW ANSWER

"A document loader is used to load data from different sources into a format that LangChain can process."

Q042. What is Text Splitting?

Large documents are usually divided into smaller pieces called chunks.

EXAMPLE

```
Large PDF
  ↓
Text
  ↓
Chunk 1
Chunk 2
Chunk 3
Chunk 4
```

Q043. Why?

Because sending the entire document to the LLM for every question is inefficient.

INTERVIEW ANSWER

"Text splitting divides large documents into smaller chunks so they can be embedded, stored and retrieved efficiently."

Q044. What is an Embedding?

An embedding converts text into numbers representing its meaning.

EXAMPLE

```
"Python is a programming language"
  ↓
Embedding Model
  ↓
[0.12, -0.45, 0.78, ...]
```

Similar meanings produce vectors that are close to each other.

Embeddings are commonly used with vector databases.

Q045. What is a Vector Store?

A Vector Store stores embeddings.

EXAMPLE

```
Documents
  ↓
Chunks
  ↓
Embedding Model
  ↓
Vectors
  ↓
Vector Store
```

EXAMPLES

- Chroma
- FAISS
- Pinecone
- Weaviate
- Milvus

When a user asks a question:

```
Question
  ↓
Embedding
  ↓
Vector Search
  ↓
Relevant Chunks
```

Q046. What is Memory in LangChain?

Memory allows an application to maintain information from previous interactions.

EXAMPLE

```
User: My name is Ravi.
Agent: Nice to meet you, Ravi.
User: What is my name?
Agent: Your name is Ravi.
```

The application needs some form of conversation history or persistent memory to do this.

Important interview point

Don't simply say:

"LangChain automatically remembers everything."

Instead say:

"Memory is implemented by maintaining and retrieving conversation or application state. The exact mechanism depends on the application architecture."

Q047. How do you connect a database with LangChain?

There are several approaches.

For example:

```
User
  ↓
Agent
  ↓
Database Tool
  ↓
SQL Database
  ↓
Result
  ↓
Agent
```

The database operation can be exposed as a tool.

For example:

```
get_customer()
get_orders()
get_sales()
```

The Agent chooses the appropriate tool.

Q048. How do you connect an API with LangChain?

You can wrap an API call as a tool.

For example:

```
Agent
  ↓
Weather Tool
  ↓
Weather API
  ↓
JSON Response
  ↓
Agent
```

The Agent can then use the API when required.

INTERVIEW ANSWER

"I normally expose external API functionality through a tool interface so the Agent can call it with validated arguments."

Q049. What is Structured Output?

Normally an LLM may return:

```
The customer is eligible for the loan.
```

But sometimes the application needs structured data:

```
{
  "customer": "Ravi",
  "eligible": true,
  "score": 780
}
```

Structured output makes it easier for application code to consume the model response.

INTERVIEW ANSWER

"Structured output forces or guides the model to return data in a predefined schema, which makes downstream processing more reliable."

Q050. How do you handle errors in LangChain?

An experienced answer should mention:

- Try/catch
- Retries
- Timeouts
- Validation
- Fallbacks
- Logging
- Error messages
- Tool error handling

EXAMPLE

```
Agent
↓
Tool
↓
ERROR
↓
Retry
↓
Success
```

If retry fails:

```
Agent
↓
Fallback
↓
Human / Error Response
```

INTERVIEW ANSWER

"I don't allow tool or model failures to break the complete workflow. I use validation, retries, timeouts, fallbacks and proper logging depending on the failure type."

Q051. LangChain vs LangGraph?

Very important.

LangChain

Useful for:

- LLM integrations
- Prompts
- Tools
- Retrievers
- Agents
- RAG
- Components

LangGraph

Useful when we need more control over:

- State
- Workflow
- Loops
- Conditional routing
- Human-in-the-loop
- Multi-Agent systems
- Long-running workflows

Simple comparison

```
LangChain
↓
Building Blocks
LangGraph
↓
Controlled Agent Workflow
```

Experienced answer

"I use LangChain for the building blocks such as models, prompts, tools and retrievers. When I need explicit control over stateful and complex Agent workflows, I use LangGraph."

Q052. Is LangChain mandatory for building an AI Agent?

No.

You can build an Agent directly using:

```
Python
+
LLM API
+
Tool Calling
+
Your own orchestration
```

But LangChain provides many reusable components and integrations.

INTERVIEW ANSWER

"No. LangChain is not mandatory. We can build an Agent directly using an LLM provider's APIs and our own orchestration code. LangChain helps reduce development effort and provides reusable abstractions."

This is a good experienced answer because you are not treating a framework as mandatory.

Q053. When would you NOT use LangChain?

This is another question that can make your answer sound experienced.

"If the application only requires a simple LLM call with a small amount of application logic, I may not use LangChain. Adding a framework is not always necessary. I use it when its abstractions and integrations provide real value."

EXAMPLE

```
User
↓
OpenAI API
↓
Answer
```


For something this simple, direct API usage may be enough.

Q054. How would you make a LangChain application production-ready?

A strong answer:

"I separate configuration from code, secure API keys, validate inputs and outputs, handle tool failures, add retries and timeouts, implement logging and monitoring, control Agent iterations, evaluate the application with real test cases, and deploy it using containers and CI/CD."

ARCHITECTURE



Q055. How would you explain LangChain in an interview in 30 seconds?

Memorize this:

"LangChain is a framework for building applications around LLMs. It provides components for models, prompts, tools, retrievers, RAG and Agents. I use it to connect an LLM with external data and services. For simple applications, I can use chains or runnable pipelines. For dynamic tool-based applications, I can use Agents, and for complex stateful workflows I prefer LangGraph."


1. Real-Time Project Question — Interviewer:

"You need to build a customer support Agent. How would you use LangChain?"

Strong simple answer:

"First, I would connect the application to an LLM. Then I would create tools for operations such as checking customer details, order status and refund status. I would use RAG if the Agent needs company documentation or policies. The Agent can decide which tool to use based on the user's request. I would add validation, authentication, logging and error handling. If the workflow becomes complex or requires multiple steps and state management, I would use LangGraph."

ARCHITECTURE




1. Most Important LangChain Questions to Remember — Before an interview, make sure you can explain these without memorizing textbook definitions:

Topic	What you should be able to explain
LangChain	Why it is used
Prompt	Instructions to the LLM
Prompt Template	Reusable dynamic prompt
Chain	Fixed sequence of operations
LCEL	Connect LangChain components
Runnable	Executable/composable component
Tool	External capability
Agent	Dynamic decision-making
Retriever	Finds relevant information
RAG	Uses external knowledge with LLM
Loader	Loads documents/data
Chunking	Splits large documents
Embedding	Converts text to vectors
Vector Store	Stores/searches vectors
Memory	Maintains useful previous information
Structured Output	Returns data in a defined format
Error Handling	Retry, timeout, fallback, validation
LangChain vs LangGraph	Building blocks vs controlled stateful workflows

The key mental model

Don't memorize LangChain as a list of classes.

Think about it like this:



Q056. LangGraph — Interview Questions & Answers

Q057. What is LangGraph?

SIMPLE ANSWER

LangGraph is a framework for building stateful and controllable AI Agent workflows.

It represents the application as a graph.

A graph contains:

- Nodes → perform work
- Edges → decide where to go next
- State → stores information during execution

Simple architecture

```

graph TD
    User --> START
    START --> Node1[Node 1]
    Node1 --> Node2[Node 2]
    Node2 --> Node3[Node 3]
    Node3 --> END
  
```

The important advantage is that we can control how the Agent moves between different steps.

Experienced interview answer

"LangGraph is used to build stateful Agent workflows. It represents the workflow as nodes and edges, with a shared state that moves through the graph. I use it when I need control over multi-step execution, conditional routing, loops, retries, human approval or multi-agent workflows."

Q058. Why do we need LangGraph?

Suppose we have a simple application:

```
Question
↓
LLM
↓
Answer
```

We don't need LangGraph.

But imagine:

```
Question
↓
Search
↓
Check Result
↓
If good → Answer
If bad → Search Again
↓
Human Approval
↓
Final Answer
```

Now the workflow is more complex.

LangGraph helps us control this workflow.

INTERVIEW ANSWER

"I use LangGraph when the Agent workflow becomes complex and I need explicit control over state, routing, loops and execution."

Q059. What is a Graph in LangGraph?

A graph represents the flow of execution.

For example:

```
START
  ↓
Research
  ↓
Analyze
  ↓
Generate
  ↓
END
```

Here:

- Research = Node
- Analyze = Node
- Generate = Node
- Arrows = Edges

Q060. What is a Node?

A Node performs a task.

For example:

```
Research Node
  ↓
gets information
```

Another:

```
Validation Node
  ↓
checks the result
```

Another:

```
Answer Node
  ↓
generates final response
```

Simple definition

A node is a function that performs one step of the workflow.

EXAMPLE

```
def research(state):
    # perform research
    return state
```

Q061. What is an Edge?

An Edge defines where the workflow should go next.

EXAMPLE

```

Research
  ↓
Analysis
  
```

The edge connects:

```

Research → Analysis
  
```

Simple definition

An edge connects one node to another.

Q062. What is a Conditional Edge?

This is very important.

A normal edge says:

```

A → B
  
```

A conditional edge can say:

```

      A
     / \
    /   \
   /     \
  Good   Bad
   ↓     ↓
   B     C
  
```

The next node depends on a condition.

EXAMPLE

Suppose an Agent searches documents.

```

Search
  ↓
Is result relevant?
  
```

☐ ☐

```

  Yes No
   ↓ ↓
Answer Search Again
  
```

This is a conditional route.

INTERVIEW ANSWER

"A conditional edge allows the workflow to dynamically choose the next node based on the current state or result."

Q063. What is State in LangGraph?

State is one of the most important LangGraph concepts.


State contains the information shared between nodes.

EXAMPLE

```
state = {  
    "question": "...",  
    "documents": [],  
    "answer": "",  
    "status": ""  
}
```

As the workflow executes, nodes read and update this state.

Flow



INTERVIEW ANSWER

"State is the shared data carried through the graph. Each node can read the current state and return updates to it."

Q064. Why is State important?

Without state, nodes would have difficulty sharing information.

EXAMPLE

```
Node 1:  
Find customer  
Node 2:  
Analyze customer  
Node 3:  
Generate response
```

Node 2 needs the result from Node 1.

State carries that information.

```
Node 1
  ↓
State:
customer = Ravi
  ↓
Node 2
  ↓
State:
customer + analysis
  ↓
Node 3
```

INTERVIEW LINE

"State allows different parts of the workflow to share information without tightly coupling the nodes."

Q065. What are START and END?

LangGraph workflows normally have an entry point and an ending point.

```
START
  ↓
Node
  ↓
Node
  ↓
END
```

START

Where execution begins.

END

Where execution finishes.

INTERVIEW ANSWER

"START represents the beginning of graph execution and END represents successful termination of the workflow."

Q066. How does a LangGraph workflow execute?

Suppose we create:

```
START
↓
Question
↓
Research
↓
Validate
↓
Answer
↓
END
```

Execution happens approximately like this:

1. Start
2. Create initial state
3. Execute Question node
4. Update state
5. Execute Research node
6. Update state
7. Execute Validate node
8. Decide next step
9. Execute Answer
10. END

The state moves through the workflow.

Q067. What is the difference between LangChain and LangGraph?

LangChain

Provides building blocks:

```
LLM
Prompt
Tool
Retriever
Agent
Parser
```

LangGraph

Provides controlled workflow orchestration:

```
State
↓
Node
↓
Conditional Edge
↓
Loop
↓
Human Approval
↓
END
```

Best interview answer

"LangChain provides the building blocks for LLM applications, while LangGraph provides a graph-based orchestration layer for stateful and complex workflows. I can use LangChain components inside LangGraph nodes."

That's an important point.

They are not necessarily competitors.

Q068. Can LangGraph work without LangChain?

Yes.

LangGraph can be used with different model/tool integrations.

But LangChain components are commonly used inside LangGraph applications.

For example:

```
LangGraph
↓
Node
↓
LangChain Model
↓
OpenAI
```


INTERVIEW ANSWER

"LangGraph does not mean every component must come from LangChain. But LangChain integrations are commonly used inside LangGraph workflows."


Q069. What is a LangGraph Agent?

A LangGraph Agent is an Agent workflow represented using a graph.

For example:



□ □




The graph gives you more control over this execution.


Q070. What is a loop in LangGraph?

A loop means the workflow can return to a previous node.

EXAMPLE



□ □



If validation fails:

Validate → Generate

The Agent tries again.

INTERVIEW ANSWER

"LangGraph makes loops explicit. For example, a validation node can route back to a generation node when the output does not meet the required criteria."

Q071. Why are loops useful in Agentic AI?

Agents often need to:

- Try again
- Search again
- Correct an answer
- Validate output
- Retry a failed operation

EXAMPLE

```
Search
↓
Evaluate
↓
Not enough information
↓
Search Again
```

This is difficult to control cleanly in a simple linear chain.

LangGraph makes it explicit.

Q072. What is Human-in-the-Loop?

Human-in-the-loop means the Agent pauses and asks a human before continuing.

EXAMPLE

```
Agent
↓
Generate Refund
↓
Human Approval
↓
Approved?
```

□ □

```
Yes No
↓ ↓
Process Stop
```

Example use cases

- Financial transactions
- Sending emails
- Deleting data
- Refunds
- Production deployments

- Important business decisions

INTERVIEW ANSWER

"For sensitive actions, I don't allow the Agent to act completely autonomously. I add a human approval step before the action is executed."

This is a very strong production answer.

Q073. What is Checkpointing?

Checkpointing means saving the state of the workflow at different points.

EXAMPLE

```
Node 1
  ↓
Save State
  ↓
Node 2
  ↓
Save State
  ↓
Node 3
```

If something fails, the application can potentially resume from a saved state instead of starting from the beginning.

INTERVIEW ANSWER

"Checkpointing allows the workflow state to be persisted so long-running or interrupted workflows can resume from a previous point."

Q074. Why is persistence important?

Imagine a workflow takes 10 minutes:

```
Step 1
Step 2
Step 3
Step 4
Step 5
Step 6
```

If the application crashes at Step 5, without persistence we may need to start again.

With checkpointing:

```
graph TD
    A[Saved State] --> B[Resume]
    B --> C[Step 5]
```

This is especially useful for:

- Long-running Agents
- Human approval
- Multi-step workflows
- Production applications

Q075. How do you handle errors in LangGraph?

I normally use:

- Error handling
- Retry logic
- Timeouts
- Validation
- Fallback nodes
- Conditional routing
- Logging

EXAMPLE

```
graph TD
    A[Tool] --> B[Error]
    B --> C[Retry]
    C --> D[Success]
```

If retry fails:

```
graph TD
    A[Retry] --> B[Failed]
    B --> C[Fallback]
    C --> D[Human / Error Response]
```

INTERVIEW ANSWER

"I handle failures at the workflow level. Depending on the failure, I can retry the node, route to a fallback node, stop the workflow, or request human intervention."

Q076. How do you prevent infinite loops?

Suppose:

```
Generate
↓
Validate
↓
Failed
↓
Generate
↓
Validate
↓
Failed
↓
...
```

This can continue forever.

So we maintain something like:

```
retry_count = 0
```

Then:

```
retry_count < 3
```

If not:

```
END / ERROR
```

Other controls:


- Maximum iterations
- Timeout
- Retry limit
- Validation
- Explicit termination conditions

Experienced answer

"I always put execution limits on Agent loops. For example, I can maintain an iteration count and stop after a defined number of attempts."

Q077. How do you build a Multi-Agent system using LangGraph?

This is a very important interview question.

EXAMPLE


Each Agent can be represented as a node or subgraph.

INTERVIEW ANSWER

"I can model each specialized Agent as a node or subgraph and use a supervisor or routing node to decide which Agent should execute next. The shared state carries the information between them."

Q078. What is a Supervisor Agent?

A Supervisor Agent controls other Agents.

EXAMPLE

The Supervisor decides:

"This task requires the Research Agent."

After that:

"Now I need the SQL Agent."

INTERVIEW ANSWER

"A Supervisor Agent acts as an orchestrator. It decides which specialized Agent should handle the next part of the task."

Q079. What is a Subgraph?

A subgraph is basically a graph inside another graph.

For example:

```
graph TD
    A[Main Graph] --> B[Research Subgraph]
    B --> C[Analysis Subgraph]
    C --> D[Final Response]
```

The research part may itself contain:

```
graph TD
    A[Search] --> B[Validate]
    B --> C[Retry]
    C --> D[Summarize]
```

This helps keep complex applications organized.

Q080. Why is LangGraph useful for production?

Because production Agents need control.

You may need:

- State management
- Conditional routing
- Loops
- Retries
- Timeouts
- Human approval
- Persistence
- Error handling
- Observability

A simple Agent loop can become difficult to manage as complexity increases.


INTERVIEW ANSWER

"For production Agent systems, I value predictability and control. LangGraph gives me explicit workflow structure, state management and controlled transitions, which makes complex Agent behavior easier to manage."

Q081. How would you design a Customer Support Agent using LangGraph?

This is a great project-based interview question.

ARCHITECTURE



EXAMPLE

User:

"Where is my order?"

Graph:

```
START
↓
Classify
↓
Order Query
↓
Order Database Tool
↓
Get Order Status
↓
Generate Response
↓
END
```

User:

"What is your refund policy?"

Graph:

```

START
↓
Classify
↓
FAQ
↓
RAG
↓
Retrieve Refund Policy
↓
Generate Response
↓
END

```

This demonstrates conditional routing.

Q082. How would you design an Agentic RAG system using LangGraph?

This is especially useful for your interviews.

```

      User
      ↓
    Question
      ↓
  Query Analysis
      ↓
    Retrieve Docs
      ↓
    Check Quality
    /  \
   /    \
Good Bad
  ↓  ↓
Generate Rewrite Query
  ↓  ↓
Answer Retrieve Again
  ↓
    END

```

This is much better than simply saying:

"I use LangChain for RAG."

You are showing that you understand Agentic RAG workflow design.

Q083. LangChain Agent vs LangGraph Agent?

Traditional Agent approach

```

LLM
↓
Choose Tool
↓
Tool
↓
LLM
↓
Choose Tool
↓
...

```

LangGraph approach

```

START
↓
Agent Node
↓
Tool Node
↓
Validation Node
↓
Conditional Edge

```

□ □

```

Retry Answer
↓ ↓
Agent END

```

Experienced answer

"A basic Agent loop can work well for simple tasks. When I need explicit control over state, routing, retries, loops or human approval, I prefer implementing the workflow using LangGraph."

Q084. How do you debug a LangGraph application?

I would check:

1. State — What is the current state?
2. Node — Which node executed?
3. Tool — Which tool was called?
4. Input — What arguments were passed?
5. Output — What did the tool return?
6. Routing — Why did the graph choose the next node?
7. LLM — What response did the model produce?
8. Performance — How many calls were made?

INTERVIEW ANSWER

"I trace the execution path and inspect state changes, node outputs, tool calls, routing decisions, LLM responses, latency and errors."

Q085. How do you control Agent behavior in LangGraph?

I don't give the Agent unlimited freedom.

I can control:

```
Tools
↓
Permissions
↓
State
↓
Conditions
↓
Maximum iterations
↓
Human Approval
↓
END
```

Experienced answer

"I use the graph structure itself as a control mechanism. The Agent can only move through the nodes and tools that I expose, and I can add conditions, validation and human approval around sensitive operations."

Q086. What are the advantages of LangGraph?

Main advantages

1. State management — Easy to maintain workflow state.
2. Explicit workflow — You can clearly see the execution flow.
3. Conditional routing — Different paths can be taken based on results.
4. Loops — Useful for retry and validation.
5. Human-in-the-loop — Easy to introduce approval points.
6. Multi-Agent — Useful for coordinating multiple Agents.
7. Persistence — Useful for long-running workflows.
8. Control — Better control than an uncontrolled Agent loop.

Q087. What are the disadvantages/challenges?

LangGraph also adds complexity.

You need to understand:

- State design
- Graph design
- Nodes
- Edges
- Routing
- Persistence
- Error handling

For a simple application, this may be unnecessary.

Experienced answer

"The main trade-off is complexity. I wouldn't introduce LangGraph for a simple one-step LLM application. I use it when the workflow complexity justifies the additional orchestration layer."

Q088. When should you NOT use LangGraph?

If the application is simply:

```
User
↓
Prompt
↓
LLM
↓
Answer
```

LangGraph is probably unnecessary.

Even:

```
User
↓
Retriever
↓
LLM
↓
Answer
```

may not require a complex graph.

INTERVIEW ANSWER

"I don't use LangGraph just because it is an Agentic AI project. If the workflow is simple and deterministic, a normal pipeline is usually easier to maintain."

Q089. Explain LangGraph in 30 Seconds

This is worth memorizing:

"LangGraph is a framework for building stateful Agent workflows. It represents the workflow as a graph containing nodes and edges, while state carries information between the nodes. I use it when I need complex orchestration such as conditional routing, loops, retries, human-in-the-loop, persistence or multi-agent workflows. LangChain components such as LLMs, tools and retrievers can be used inside the LangGraph nodes."

Q090. Explain LangGraph using a Real Project


Interviewer:

"Tell me how you used LangGraph in a project."

You can answer:

"I used LangGraph to build an Agentic RAG workflow. The user first sends a question. The graph analyzes the query and retrieves relevant documents. A validation node checks whether the retrieved information is sufficient. If the information is not sufficient, the workflow rewrites the query and performs retrieval again. If the information is sufficient, the LLM generates the final response. I used state to maintain the question, retrieved documents and intermediate results, and conditional edges to control the workflow."

ARCHITECTURE



That answer demonstrates:


Agent + RAG + LangGraph + State + Conditional Routing + Loop

—which is exactly the kind of combination interviewers often want to hear.

Final LangGraph Revision Sheet

Question	Simple answer
What is LangGraph?	Framework for stateful Agent workflows
Graph?	Represents workflow
Node?	Performs a task
Edge?	Connects nodes
Conditional Edge?	Chooses next node based on a condition
State?	Shared information during execution
START?	Beginning of workflow
END?	End of workflow
Loop?	Goes back to an earlier node
Checkpoint?	Saves workflow state
Human-in-loop?	Human approval before continuing
Subgraph?	Graph inside another graph
Supervisor?	Controls other Agents
Multi-Agent?	Multiple specialized Agents working together
LangChain?	LLM application building blocks
LangGraph?	Stateful workflow orchestration
Production benefit?	Control, state, routing, retries and persistence
Main challenge?	Adds complexity
Infinite loop prevention?	Iteration limits + conditions + timeouts
Error handling?	Retry + fallback + routing + logging

The one mental model to remember




Q091. Single Agent — Interview Questions & Answers

Q092. What is a Single Agent?

SIMPLE ANSWER

A Single Agent is one AI Agent that uses an LLM, tools, memory and reasoning to complete a task.

EXAMPLE



One Agent is responsible for the complete task.

Q093. How is a Single Agent different from a normal chatbot?

A normal chatbot may simply do:

```
User
↓
LLM
↓
Answer
```

A Single Agent can do:


```
User
↓
Agent
↓
Understand Goal
↓
Decide Action
↓
Use Tool
↓
Observe Result
↓
Decide Next Action
↓
Final Answer
```

INTERVIEW ANSWER

"A normal chatbot mainly generates responses. A Single Agent can dynamically decide what actions are required and use tools to complete the task."

Q094. What are the main components of a Single Agent?

A typical Single Agent contains:



Main components:

- LLM – reasoning engine

- Prompt – Agent instructions
- Tools – external actions
- Memory/State – information about the task
- Orchestration – controls execution
- Guardrails – controls what the Agent can do

Q095. How does a Single Agent work?

Suppose the user asks:

"Check my order status."

The Agent may work like this:

```
User
↓
Agent
↓
LLM
↓
Understand request
↓
Select Order Tool
↓
Database/API
↓
Order Status
↓
LLM
↓
Final Answer
```

The important point is:

The Agent decides which action is needed.

Q096. How does the Agent decide which tool to use?

The available tools are provided to the LLM with their descriptions.

For example:

```
Tool 1: get_order_status(order_id)
Tool 2: cancel_order(order_id)
Tool 3: get_customer_details(customer_id)
```

User:

"Where is my order?"

The LLM identifies that `get_order_status()` is the appropriate tool.

Then the Agent executes that tool.

INTERVIEW ANSWER


"The LLM uses the tool descriptions and the user's request to select the appropriate tool. The application then executes the selected tool and returns the result to the Agent."

Q097. Can a Single Agent use multiple tools?

Yes.

One Agent can have many tools.

EXAMPLE



For example, a customer-support Agent could have:

- Customer lookup
- Order lookup
- Refund
- Product search
- Email

The same Agent can choose the required tool.

Q098. Can a Single Agent use multiple tools in one task?

Yes.

For example:

"Find my order and email me the latest status."

The Agent might do:

```
User
↓
Agent
↓
Order Tool
↓
Get Status
↓
Email Tool
↓
Send Email
↓
Final Answer
```

This is one Agent using two tools.

Q099. What is the Agent Loop?

A Single Agent commonly follows this pattern:

```
Understand
↓
Reason
↓
Choose Tool
↓
Execute Tool
↓
Observe Result
↓
Reason Again
↓
Choose Next Action
↓
...
↓
Final Answer
```

EXAMPLE

```
User
↓
LLM
↓
Search Tool
↓
Search Result
↓
LLM
↓
Database Tool
↓
Database Result
↓
LLM
↓
Final Answer
```

This is the Agent loop.

Q100. Does the Agent execute tools itself?

This is an important technical distinction.

The LLM decides which tool should be called, but the application or Agent runtime executes the tool.

Conceptually:

```
LLM
↓
"Call get_order_status"
↓
Agent Runtime
↓
Execute Function
↓
Return Result
↓
LLM
```

Experienced answer

"The model produces a structured tool call, but the actual tool execution is handled by the application or Agent runtime."

Q101. What is the role of the LLM in a Single Agent?

The LLM acts as the reasoning and decision-making component.

It can help decide:

- What does the user want?
- Do I need a tool?
- Which tool should I use?
- What arguments should I provide?
- Do I need another step?
- Is the task complete?

SIMPLE ANSWER

"The LLM is the brain of the Agent, while tools allow the Agent to interact with the external world."

Q102. What is the role of tools?

Tools allow the Agent to do things outside the LLM.

EXAMPLES

```
LLM
↓
Tool
├─ Database
├─ API
├─ Search
├─ Calculator
├─ File System
└─ Email
```

Without tools, the Agent has limited ability to interact with external systems.

INTERVIEW LINE

"Tools give the Agent access to capabilities that the LLM itself does not have."

Q103. What is the role of memory?

Memory allows the Agent to retain useful information.

EXAMPLE

```
User: My customer ID is 1001.
Agent: Okay.
Later:
User: Check my order.
Agent → Uses customer ID 1001
```

Memory can be:

- Conversation history
- Database
- Vector store
- Key-value store

- Application state

Important:

Don't say:

"Every Agent automatically has permanent memory."

Instead:

"Memory is an additional capability that we implement when the application requires it."

Q104. What is State in a Single Agent?

State contains information about the current task.

EXAMPLE

```
state = {
  "question": "Where is my order?",
  "customer_id": "1001",
  "order_id": "ORD123",
  "tool_result": None,
  "final_answer": None
}
```

During execution, the Agent updates this information.

INTERVIEW ANSWER

"State represents the current context of the task and helps the Agent keep track of what has already happened."


Q105. When should you use a Single Agent?

Use a Single Agent when:

- One Agent can handle the complete problem.
- There are multiple tools.
- The task requires dynamic decisions.
- The workflow isn't extremely complex.
- There is no strong need to separate responsibilities.

EXAMPLE

Customer Support Agent:



One Agent can manage all these operations.

Q106. Why not use Multi-Agent for everything?

This is a very good experienced-level question.

More Agents don't automatically mean a better system.

Multi-Agent introduces:

- More complexity
- More LLM calls
- More communication
- More latency
- More cost
- More failure points
- More debugging effort


Strong interview answer

"I don't introduce Multi-Agent architecture unless there is a clear need for specialization or independent responsibilities. If one Agent with well-defined tools can solve the problem, I prefer Single Agent because it is simpler to develop, test and maintain."

This is a strong architectural answer.


Q107. Single Agent vs Multi-Agent

Single Agent



One Agent controls everything.

Multi-Agent



Multiple specialized Agents.

INTERVIEW ANSWER

"In a Single-Agent system, one Agent handles the task using multiple tools. In a Multi-Agent system, different Agents are responsible for different tasks or responsibilities."

Q108. What are the advantages of a Single Agent?

1. Simple architecture — Fewer components.
2. Easier debugging — One Agent's execution is easier to trace.
3. Lower cost — Usually fewer LLM calls than a comparable Multi-Agent design.
4. Lower latency — Less Agent-to-Agent communication.
5. Easier deployment — Fewer services/components.
6. Easier maintenance — Less orchestration code.


INTERVIEW ANSWER

"The main benefits are simplicity, lower operational overhead, easier debugging and generally lower latency and cost."

Q109. What are the limitations of a Single Agent?

A Single Agent can become difficult when it has too many responsibilities.

EXAMPLE



Now the Agent has too many tools and responsibilities.

Problems may include:

- Complex prompts
- Tool selection errors
- Large context
- Difficult debugging
- Poor separation of responsibilities

Experienced answer

"A Single Agent can become difficult to manage when it has too many tools or unrelated responsibilities. At that point, I consider splitting responsibilities into specialized Agents or deterministic workflows."

Q110. What is Tool Selection?

Tool selection means:

The Agent decides which tool is appropriate for the current task.

EXAMPLE

```
User:
"Show my account balance."
Agent
↓
Tool Selection
↓
get_balance()
```

Another request:

```
User:
"Cancel my order."
Agent
↓
Tool Selection
↓
cancel_order()
```

Q111. How do you prevent an Agent from using the wrong tool?

Use:

1. Clear tool descriptions — Bad:

```
customer_tool()
```

Better:

```
get_customer_details(customer_id)
```

1. Strong input validation — Check arguments before execution.
2. Tool permissions — Don't expose unnecessary tools.
3. Guardrails — Validate dangerous operations.
4. Human approval — For sensitive operations.

INTERVIEW ANSWER

"I reduce incorrect tool usage by exposing only necessary tools, giving clear descriptions, validating arguments and adding authorization and approval checks for sensitive operations."

Q112. How do you handle a dangerous tool?

Suppose the Agent has:

```
delete_customer()
```

You should not allow the Agent to execute it freely.

Instead:

```
Agent
↓
Delete Request
↓
Permission Check
↓
Human Approval
↓
Delete
```

INTERVIEW ANSWER

"I treat tools as privileged operations. The Agent should not automatically have unrestricted access to sensitive tools. I apply authorization, validation and human approval where required."

Q113. How do you prevent infinite loops in a Single Agent?

Use:

```
Maximum iterations
Maximum tool calls
Timeout
Retry limit
Termination condition
```

EXAMPLE

```
max_iterations = 10
```

Then:

```
Agent
↓
Tool
↓
Agent
↓
Tool
↓
...
↓
10 iterations
↓
STOP
```

INTERVIEW ANSWER

"I always define execution limits. I use maximum iterations, timeouts and explicit termination conditions so the Agent cannot run indefinitely."

Q114. How do you handle tool failures?

EXAMPLE

```
Agent
↓
Payment API
↓
ERROR
```

The Agent can:

```
Retry
↓
Success
```

Or:

```
Retry
↓
Still Failed
↓
Fallback
↓
Human / Error Response
```

INTERVIEW ANSWER

"I classify tool failures and handle them using retries, timeouts, fallbacks and proper error messages. For critical operations, I avoid blindly retrying actions that may cause duplicate transactions."

That last sentence sounds particularly experienced.

Q115. How do you handle hallucinations in a Single Agent?

Don't depend on the LLM for information that should come from a trusted system.

Instead:

```
Question
↓
Agent
↓
Trusted Tool
↓
Database/API
↓
Actual Data
↓
LLM
↓
Answer
```

For company knowledge:

```
Question
↓
Agent
↓
RAG
↓
Company Documents
↓
LLM
```

INTERVIEW ANSWER

"For business-critical information, I ground the Agent using trusted APIs, databases or RAG instead of relying only on the model's internal knowledge."

Q116. How do you evaluate a Single Agent?

Don't evaluate only the final answer.

Check:

Task completion

Did it complete the task?

Tool selection

Did it select the correct tool?

Arguments

Were the tool arguments correct?

Accuracy

Was the final answer correct?

Efficiency

How many LLM calls?

Q117. How many tool calls?

Performance

How long did it take?

Cost

How many tokens/API calls?

Safety

Did it perform unauthorized actions?

INTERVIEW ANSWER

"I evaluate both the final output and the Agent's execution. I check task completion, tool selection, tool arguments, accuracy, latency, token usage, cost and safety."

Q118. How do you monitor a Single Agent in production?

I would monitor:

```
Requests
↓
LLM Calls
↓
Tool Calls
↓
Errors
↓
Latency
↓
Token Usage
↓
Cost
↓
Success Rate
```

I would also maintain logs and traces.

INTERVIEW ANSWER

"In production, I want visibility into the complete Agent execution path, including LLM calls, tool calls, errors, latency, token usage and task success."

Q119. How do you secure a Single Agent?

Important areas:

Authentication

Who is using the Agent?

Authorization

What is the user allowed to do?

Tool permissions

Which tools can the Agent access?

Secrets

API keys should not be inside prompts or source code.

Input validation

Validate user and tool inputs.

Output validation

Validate important outputs.

Human approval

For high-risk actions.


INTERVIEW ANSWER

"I treat the Agent as an application with access to privileged tools. So I apply authentication, authorization, tool-level permissions, secret management, input validation and human approval for sensitive actions."

Q120. Can a Single Agent use RAG?

Yes.

EXAMPLE



The Agent decides whether it needs:

- Document retrieval
- Database
- API

- Search

EXAMPLE

User:

"What is our refund policy?"

Agent → RAG

User:

"What is my refund status?"

Agent → Customer/Refund API


Same Agent, different tools.

Q121. Can a Single Agent use MCP?

Yes.

MCP can expose tools that the Agent can use.

Conceptually:



The Agent can access capabilities exposed through MCP.

Q122. Can a Single Agent be built using LangGraph?

Yes.

For example:

```
START
↓
Agent
↓
Tool
↓
Agent
↓
Tool
↓
END
```

LangGraph can provide:

- State
- Tool execution
- Conditional routing
- Loops
- Retry
- Human approval
- Persistence

INTERVIEW ANSWER

"Yes. A Single Agent can be implemented as a LangGraph workflow. LangGraph is especially useful when I need explicit control over the Agent's state and execution."

Q123. Real Project Interview Question


Interviewer:

"Design a Single Agent for an e-commerce application."

Strong answer:

"I would create one customer-support Agent with access to a limited set of tools. The tools could include customer lookup, order status, product search and refund status. The LLM would decide which tool is required based on the user's request. I would maintain the conversation and task state, validate tool inputs, enforce authorization and add human approval for sensitive actions such as refunds. If the workflow becomes too complex, I would consider splitting responsibilities."

ARCHITECTURE



Q124. "Why did you choose Single Agent?"

This is a very common architecture question.

Don't say:

"Because it is simple."

Give a technical reason:

"I chose Single Agent because the use case has one main responsibility and the required capabilities can be exposed as tools. Using multiple Agents would add communication and orchestration overhead without providing enough benefit. So I kept one Agent and controlled its tools, permissions and execution."

Excellent interview answer.

Q125. "When would you move from Single Agent to Multi-Agent?"

Say:

"I would consider Multi-Agent when the problem has clearly separated responsibilities, different expertise, independent workflows or when one Agent becomes difficult to manage because of too many tools and instructions."

EXAMPLE

```

Single Agent
  ↓
Too many responsibilities
  ↓
Split
  ↓
Supervisor
 / | \
Research Coding Testing

```

Q126. Explain Single Agent in 30 Seconds

Memorize this:


"A Single Agent is one AI Agent that uses an LLM as its reasoning engine and can use multiple tools to complete a task. The Agent receives the goal, decides what action is needed, calls the appropriate tool, observes the result and continues until the task is complete. I prefer Single Agent when one Agent can handle the complete responsibility because it is simpler to develop, test, monitor and maintain."

Q127. Single Agent — Quick Revision

Question	Simple Answer
What is Single Agent?	One Agent handling a task using LLM + tools
Brain of Agent?	LLM
What gives Agent capabilities?	Tools
Can it use multiple tools?	Yes
Can it use RAG?	Yes
Can it use MCP?	Yes
Can it use memory?	Yes
Can it use LangGraph?	Yes
Who decides the tool?	Usually the LLM
Who executes the tool?	Application/Agent runtime
What is Agent loop?	Decide → Act → Observe → Repeat
Why Single Agent?	Simpler architecture
Main benefit?	Easier development and maintenance
Main limitation?	Can become complex with too many responsibilities

Question	Simple Answer
How prevent loops?	Limits + timeout + termination conditions
How handle failures?	Retry + fallback + error handling
How reduce hallucination?	RAG + trusted APIs/tools
How secure it?	Authentication + authorization + tool permissions
When Multi-Agent?	When responsibilities need clear separation

One mental model




Q128. Multi-Agent Systems — Interview Questions & Answers

Q129. What is a Multi-Agent System?

SIMPLE ANSWER

A Multi-Agent system is a system where multiple AI Agents work together to complete a complex task.

Each Agent can have a different responsibility.

EXAMPLE

For example:

- Research Agent → finds information
- Coding Agent → writes code
- Testing Agent → tests the code
- Supervisor → decides which Agent should work next

Q130. Why do we need Multi-Agent Systems?

A single Agent can handle many tasks.

But sometimes the problem becomes too large for one Agent.


For example:

Build an AI application


This may involve:

Requirements
Research
Coding
Testing
Security
Deployment
Documentation

Instead of giving all responsibilities to one Agent:



we can use specialized Agents:



Experienced answer

"I use Multi-Agent architecture when the problem has clearly separated responsibilities that can be handled by specialized Agents. This can make the system easier to organize, but it also adds communication, latency and cost."

Q131. What is an Agent's role in a Multi-Agent system?

Each Agent normally has a specific responsibility.

EXAMPLE

```

Research Agent
→ Find information
Coding Agent
→ Write code
Testing Agent
→ Test code
Documentation Agent
→ Create documentation
  
```


This is called specialization.

Q132. What is Agent Specialization?

Agent specialization means:

Give each Agent a specific job.

For example:



Instead of:

One Agent does everything.

We have:

Each Agent focuses on one responsibility.


INTERVIEW ANSWER

"I prefer specialized Agents when their responsibilities are clearly different. This helps keep prompts, tools and responsibilities focused."

Q133. What is a Supervisor Agent?

A Supervisor Agent manages other Agents.

EXAMPLE



The Supervisor decides:

"Which Agent should work next?"

For example:

```
User:
Build a Python application.
Supervisor:
→ Research Agent
Research Agent:
→ Research completed.
Supervisor:
→ Coding Agent
Coding Agent:
→ Code completed.
Supervisor:
→ Testing Agent
Testing Agent:
→ Testing completed.
```

Q134. What is an Orchestrator?

An orchestrator controls the overall execution of multiple Agents.

It can be:

- An Agent
- A workflow
- A deterministic program
- A graph

EXAMPLE

```
User
↓
Orchestrator
↓
Agent 1
↓
Agent 2
↓
Agent 3
↓
Final Result
```

Important distinction

A Supervisor Agent is usually an Agent making routing decisions.

An Orchestrator is a broader term for the component controlling execution.

INTERVIEW ANSWER

"Orchestration means controlling how Agents interact and how the overall workflow progresses. The orchestrator may use an LLM for dynamic routing or deterministic rules for predictable workflows."

Q135. How do Agents communicate with each other?

Agents can communicate through:

1. Shared State

```
Agent A
↓
Shared State
↓
Agent B
```

1. Messages

```
Agent A
↓
Message
↓
Agent B
```

1. Shared Database

```
Agent A
↓
Database
↓
Agent B
```

1. API

```
Agent A
↓
API
↓
Agent B
```

In LangGraph, shared state is commonly used to pass information between nodes or Agent workflows.

Q136. What is Shared State?

Shared state is information that multiple Agents can access.

EXAMPLE

```
state = {
    "user_query": "...",
    "research_result": "...",
    "code": "...",
    "test_result": "...",
}
```

Flow:

```
Research Agent
↓
research_result
↓
Shared State
↓
Coding Agent
↓
code
↓
Shared State
↓
Testing Agent
```

INTERVIEW ANSWER

"Shared state provides a common place for Agents to exchange information and maintain the progress of the overall task."

Q137. What is Agent Delegation?

Delegation means:

One Agent gives a task to another Agent.

EXAMPLE

```

Main Agent
  ↓
"I need research."
  ↓
Research Agent
  ↓
Research Result
  ↓
Main Agent

```

Another example:

```

Supervisor
  ↓
Coding Agent
  ↓
"Write login API"

```

INTERVIEW ANSWER

"Delegation allows one Agent or orchestrator to assign a specific responsibility to a specialized Agent."

Q138. What are the common Multi-Agent architectures?

There are several patterns.

1. Supervisor pattern

```

      Supervisor
     /  |  \
    ↓  ↓  ↓
   A  B  C

```

1. Sequential pattern

```

Agent A
  ↓
Agent B
  ↓
Agent C

```

1. Hierarchical pattern

```

      Manager
     /  \
  Manager Manager
   / \ / \
  A  B C  D

```

1. Peer-to-peer

```


Agent A ↔ Agent B
  ↑   ↓
Agent C ↔ Agent D

```

1. Group/Collaborative — Multiple Agents work together around shared state or messages.

Q139. What is the Supervisor Pattern?

This is one of the most common designs.



The Supervisor:

- Receives the task.
- Understands the requirement.
- Chooses an Agent.
- Receives the result.
- Decides what to do next.
- Continues until the task is complete.

Experienced answer

"I use the Supervisor pattern when I want one central component to coordinate specialized Agents. It gives me a clear control point for routing, permissions and termination."

Q140. What is Sequential Multi-Agent Architecture?

Agents execute one after another.

EXAMPLE

```

User
↓
Research Agent
↓
Analysis Agent
↓
Writing Agent
↓
Review Agent
↓
Final Result

```


This is useful when each step depends on the previous step.

EXAMPLE

Research → Analyze → Write → Review

Q141. What is Parallel Multi-Agent Architecture?

Multiple Agents can work at the same time.

EXAMPLE

Advantage


Potentially faster because independent tasks run simultaneously.

Challenge

You need to combine and validate the results.

Q142. What is Hierarchical Multi-Agent Architecture?

There can be multiple levels of Agents.



The top-level Agent manages lower-level Agents.

This is useful for very large workflows.

Q143. What is Peer-to-Peer Agent Architecture?

There is no single supervisor.

Agents communicate directly.

Agent A

□

□ □

Agent B ↔ Agent C

□ □

□

Agent D


This can be flexible but becomes harder to control.

INTERVIEW ANSWER


"Peer-to-peer communication can work when Agents need direct collaboration, but I would be careful because unrestricted Agent-to-Agent communication can make the system difficult to debug and control."

Q144. Single Agent vs Multi-Agent

Single Agent



Multi-Agent



Single Agent

Good when:

- One responsibility
- Moderate number of tools
- Simple workflow

Multi-Agent

Useful when:

- Different specialized responsibilities

- Complex workflows
- Independent tasks
- Different tools or permissions
- Large systems

Q145. Why not simply give all tools to one Agent?

This is an excellent experienced-level question.

Imagine:

```
Agent
↓
50 Tools
↓
Research
Database
Email
Payments
Coding
Deployment
Testing
Monitoring
...
```

The Agent now has too many choices.

Potential problems:

- Tool selection becomes harder.
- Prompt becomes larger.
- Context increases.
- Debugging becomes harder.
- Security boundaries become unclear.

Instead:


```
Research Agent → Research Tools
Coding Agent → Coding Tools
Testing Agent → Testing Tools
```

INTERVIEW ANSWER

"I don't expose every tool to every Agent. I prefer least-privilege tool access and specialized Agents when the responsibilities are sufficiently different."

Q146. How do you decide whether to use Single Agent or Multi-Agent?

Use this decision process:



Strong interview answer

"I start with the simplest architecture. If one Agent with controlled tools can solve the problem reliably, I use Single Agent. I introduce Multi-Agent only when specialization, isolation or parallel execution provides a clear benefit."

Q147. What are the advantages of Multi-Agent Systems?

1. Specialization — Each Agent focuses on one task.
2. Separation of responsibility — Cleaner architecture.
3. Parallel execution — Independent tasks can run simultaneously.
4. Different tools — Each Agent can have different tools.
5. Different instructions — Each Agent can have a focused prompt.
6. Easier scaling of responsibilities — You can improve one specialized Agent without changing everything.

INTERVIEW ANSWER

"The main benefits are specialization, separation of responsibilities, controlled tool access and the ability to execute independent tasks in parallel."

Q148. What are the disadvantages?

This is very important.

Multi-Agent is not automatically better.

Problems include:

Cost

More Agents can mean more LLM calls.

Latency

Agents need to communicate.

Complexity

More components to manage.

Failure propagation

One Agent's failure may affect others.

Debugging

More difficult execution paths.

Context management

Information needs to be passed correctly.

Security

More Agents may mean more tool permissions.

INTERVIEW ANSWER

"The main trade-offs are additional complexity, latency, cost, communication overhead and more failure points."

Q149. How do you prevent Agents from duplicating work?

Use:

- Clear responsibilities
- Explicit task boundaries
- Shared state
- Task IDs
- Workflow rules
- Central orchestration

EXAMPLE

```
Research Agent
→ Only research
Coding Agent
→ Only coding
Testing Agent
→ Only testing
```

Don't give all Agents the same responsibility.

INTERVIEW ANSWER

"I define clear Agent responsibilities and maintain shared task state so Agents know what has already been completed."


Q150. How do you prevent Agent conflicts?

Suppose:

Agent A → recommends refund
Agent B → rejects refund

Q151. Who decides?

Use a clear authority model.



The Supervisor can apply the final business rule.

Experienced answer

"I avoid letting multiple Agents independently perform conflicting business actions. I use an orchestrator, validation layer or deterministic business rules for final decisions."

Q152. Should Agents directly modify databases?

They can, but for sensitive systems I prefer controlled tools.

Instead of:

Agent → Database

use:

```

Agent
↓
Validated Tool
↓
Authorization
↓
Business Logic
↓
Database
  
```

EXAMPLE

```
Agent
↓
cancel_order()
↓
Check User Permission
↓
Check Order Status
↓
Execute Cancellation
```

This gives better control.

Q153. How do you secure a Multi-Agent system?

Important principles:

Authentication

Who is calling the system?

Authorization

What can the user do?

Agent permissions

What can each Agent do?

Tool permissions

Which tools can each Agent access?

Secret management

Keep API keys outside the code.

Validation

Validate inputs and outputs.

Human approval

For sensitive operations.

Least privilege

Give each Agent only the access it needs.

Strong answer

"I follow least privilege. Each Agent gets only the tools and permissions required for its responsibility. Sensitive operations go through authorization, validation and, when necessary, human approval."

Q154. How do Agents share information?

There are several approaches:

```

Agent A
  ↓
Shared State
  ↓
Agent B

```

or:

```

Agent A
  ↓
Message
  ↓
Agent B

```

or:

```

Agent A
  ↓
Database
  ↓
Agent B

```

or:

```

Agent A
  ↓
API
  ↓
Agent B


```

The choice depends on the architecture.

Q155. How would you implement Multi-Agent using LangGraph?

This is a very important interview question.

A simple design:



Each Agent can be represented as:

- A node
- A group of nodes
- A subgraph

The shared state contains the information required by the workflow.

INTERVIEW ANSWER


"With LangGraph, I can represent the Supervisor and specialized Agents as nodes or subgraphs. Shared state carries the task information, and conditional edges determine which Agent should execute next."

Q156. Can Multi-Agent systems use MCP?

Yes.

Each Agent can have access to specific MCP servers/tools.

EXAMPLE



This is useful because tools can be exposed through a standard protocol.

Q157. Multi-Agent + MCP Example

Imagine a software development system.



Research Agent can search documentation.

Coding Agent can work with repositories.

The Supervisor coordinates them.

Q158. How do you handle failures in Multi-Agent systems?

EXAMPLE

```
Supervisor
  ↓
Research Agent
  ↓
ERROR
```

Options:

```
Retry
  ↓
Success
```

or:

```
Retry
  ↓
Still Failed
  ↓
Fallback Agent
```

or:

```
Failure
  ↓
Human
```

Important:

Don't blindly retry actions that may have side effects.

For example:

```
send_payment()
```

should not be blindly repeated.

Experienced answer

"I classify failures by whether the operation is safe to retry. For read operations, retries are usually straightforward. For side-effecting operations, I use idempotency, validation and explicit safeguards."

That is a strong experienced-level answer.

Q159. How do you evaluate a Multi-Agent system?

Don't only check the final answer.

Check:

Agent selection

Was the correct Agent selected?

Task completion

Was the task completed?

Communication

Did Agents exchange the correct information?

Tool usage

Were the right tools used?

Accuracy

Was the result correct?

Cost

How many LLM calls?

Latency

How long did the workflow take?

Failure rate

How often did Agents fail?

INTERVIEW ANSWER

"I evaluate both the final result and the coordination behavior: Agent routing, tool usage, state transitions, task completion, latency, token usage, cost and failure rate."

Q160. What is Agent-to-Agent communication?

Agent-to-Agent communication means one Agent sends information or a task to another Agent.

EXAMPLE

Research Agent

↓

"Here are the research findings."

↓

Analysis Agent

The communication can contain:

Task

Context

Result

Status

Errors

Metadata

Q161. What is Agent Handoff?

A handoff means:

One Agent transfers responsibility to another Agent.

EXAMPLE

```
Customer Agent
  ↓
"This is a billing issue."
  ↓
Billing Agent
```

The Billing Agent now handles the request.

EXAMPLE

```
Customer Agent
  ↓
Technical issue detected
  ↓
Technical Support Agent
```

Q162. What is a Shared-State Multi-Agent system?

EXAMPLE

```
      Shared State
    /  |  \
  ↓ ↓ ↓
Research Coding Testing
Agent Agent Agent
```

All Agents can read or update relevant parts of the state.

This is especially useful in LangGraph.

Q163. What is a Message-Based Multi-Agent system?

Instead of shared state, Agents exchange messages.

```
Agent A
  ↓
Message Queue
  ↓
Agent B
  ↓
Message Queue
  ↓
Agent C
```

This can be useful when Agents are independent services.

For example:

```
Kafka / Queue
```

can be used between services in a larger architecture.

Q164. What is the biggest challenge in Multi-Agent systems?

A strong answer:

"Coordination."

Because now you need to manage:

```
Who works?
  ↓
When do they work?
  ↓
What information do they receive?
  ↓
What tools can they use?
  ↓
What happens if they fail?
  ↓
Who decides the final result?
  ↓
When should the workflow stop?
```

Experienced answer

"The biggest challenge is not creating multiple Agents. It is coordinating them reliably. I need clear responsibilities, controlled communication, shared state, failure handling and a well-defined termination strategy."

Q165. Real Project Interview Question


Interviewer:

"Design a Multi-Agent system for software development."

A strong answer:

"I would use a Supervisor-based architecture. The Supervisor receives the development requirement and routes work to specialized Agents. A Research Agent can analyze requirements and documentation, a Coding Agent can implement the solution, and a Testing Agent can validate the implementation. They can share progress through state. I would restrict each Agent to the tools it needs. For sensitive actions such as production deployment, I would add approval and authorization. LangGraph can orchestrate the workflow, while MCP can provide standardized access to external tools."

ARCHITECTURE



Q166. How would you answer: "Why did you choose Multi-Agent?"

Don't say:

"Because Multi-Agent is the latest technology."

Instead:

"I chose Multi-Agent because the problem had clearly separated responsibilities. Research, implementation and testing required different tools and instructions. Separating them made the responsibilities clearer and allowed me to control the permissions of each Agent. I also considered Single Agent, but the number of responsibilities and tools made specialization useful."

That sounds much more experienced.

Q167. How would you answer: "Why didn't you use Multi-Agent?"

This is equally important.

"I evaluated Multi-Agent, but the problem didn't have enough separation of responsibilities to justify the additional complexity. A Single Agent with controlled tools was sufficient, so I chose the simpler architecture."

This demonstrates architectural judgment.

Q168. How do you reduce cost in Multi-Agent systems?

Use:

- Smaller models for simple tasks
- Strong models only for complex reasoning
- Limit Agent iterations
- Avoid unnecessary Agent handoffs
- Cache repeated results
- Run independent tasks in parallel
- Reduce unnecessary context
- Use deterministic logic where possible

Experienced answer

"I don't use an expensive model for every Agent. I match the model to the task, reduce unnecessary Agent calls, control iterations and use deterministic logic for predictable operations."

Q169. How do you reduce latency?

Use:


Parallel execution
Caching
Smaller models
Fewer LLM calls
Efficient prompts
Direct APIs for deterministic operations

EXAMPLE

Instead of:

Research A
↓
Research B
↓
Research C

If independent:



They can potentially run concurrently.

Q170. Single Agent vs Multi-Agent — Interview Table

Area	Single Agent	Multi-Agent
Agents	One	Multiple
Complexity	Lower	Higher
Tools	Shared by one Agent	Can be specialized
Responsibility	Broad	Specialized
Communication	Minimal	Required
Cost	Usually lower	Can be higher
Latency	Usually lower	Can be higher
Debugging	Easier	More difficult
Parallel work	Limited	Easier
Best for	Moderate tasks	Complex tasks
Orchestration	Simple	Important
State	Useful	Very important
Security	Simpler	More complex

Q171. Explain Multi-Agent in 30 Seconds


Memorize this:

"A Multi-Agent system uses multiple specialized AI Agents to solve a complex problem. Each Agent has a specific responsibility, such as research, coding or testing. An orchestrator or supervisor controls the workflow and decides which Agent should work next. Agents can communicate using shared state, messages, APIs or other mechanisms. I use Multi-Agent only when specialization or parallel execution provides a clear benefit, because it also increases cost, latency and system complexity."

1. One Real-World Example — Suppose the user says:

"Analyze this company's performance and prepare an investment report."

A Multi-Agent system could be:




Notice the important concept:

The Agents are not created just to increase the number of Agents.

They exist because each has a clear responsibility.

Final Multi-Agent Mental Model



Q172. MCP — Model Context Protocol

Q173. What is MCP?

MCP stands for:

Model Context Protocol

SIMPLE ANSWER

MCP is an open standard that allows AI applications to connect to external tools, data and services in a standard way.

For example, an AI Agent may need access to:

- GitHub
- Database
- Files
- Search
- CRM
- Cloud services
- APIs

Instead of creating a custom integration for every AI application, we can expose capabilities through an MCP server.

The official MCP SDK describes it as a standard for connecting AI applications to systems where data and tools live. (ModelContextProtocol)

Q174. Why was MCP introduced?

Imagine you have 10 AI applications:

```
Claude  
ChatGPT  
VS Code  
Cursor  
My AI Agent  
Company AI App  
...
```

And you have:


```
GitHub  
Database  
Slack  
Files  
Jira  
AWS  
...
```

Without a standard:

```
AI App 1 → Custom GitHub integration  
AI App 2 → Custom GitHub integration  
AI App 3 → Custom GitHub integration  
AI App 4 → Custom GitHub integration
```

This becomes difficult to maintain.

With MCP:



INTERVIEW ANSWER


"MCP provides a standard interface for connecting AI applications with external tools and data, reducing the need for separate custom integrations."

Q175. What problem does MCP solve?

The main problem is:

Standardized connectivity between AI applications and external capabilities.

For example:



The AI application does not need to know every internal implementation detail of the external system.

Q176. What is an MCP Server?

This is a very common interview question.

SIMPLE ANSWER

An MCP Server exposes tools, resources and prompts that an AI application can use.

For example, a GitHub MCP server might expose:

```
Tools:
  create_issue()
  get_issue()
  create_pull_request()
Resources:
  repository information
Prompts:
  code review prompt
```

The official SDK documentation describes MCP servers as exposing tools, resources and prompts. (ModelContextProtocol)

Q177. What is an MCP Client?

SIMPLE ANSWER

An MCP Client connects an AI application to an MCP Server.

ARCHITECTURE

```
graph TD
    User --> AI_Application[AI Application]
    AI_Application --> MCP_Client[MCP Client]
    MCP_Client --> MCP_Server[MCP Server]
    MCP_Server --> External_System[External System]
```

For example:


```
graph TD
    Agent --> MCP_Client[MCP Client]
    MCP_Client --> GitHub_MCP_Server[GitHub MCP Server]
    GitHub_MCP_Server --> GitHub
```

Q178. What is an MCP Host?

This terminology is important.

The Host is the AI application that provides the environment in which MCP connections are managed.

For example, conceptually:




The host manages MCP clients and the model/application experience around them.

Simple interview answer

"The MCP Host is the AI application that manages MCP connections. An MCP Client connects the host to a particular MCP Server."


Q179. MCP Host vs Client vs Server

Remember this:



EXAMPLE

Suppose you build an AI coding application:



Q180. What is an MCP Tool?

An MCP Tool is an action that an AI application can call.

EXAMPLES

```
search_files()
read_file()
create_issue()
get_customer()
send_email()
query_database()
```

The tool normally has a defined input schema.

For example:

```
get_customer
Input:
{
  "customer_id": "1001"
}
```

The Agent can decide to call it.

Q181. What is an MCP Resource?

A Resource provides data or information that can be accessed through MCP.

EXAMPLES

```
file
document
database information
configuration
repository information
```

Think of it simply as:

Tool = Do something

Resource = Read/access something

Q182. What is an MCP Prompt?

An MCP Prompt is a reusable prompt or prompt template exposed by an MCP server.

For example:

```
code_review_prompt
```

could provide a standardized code-review instruction.

Simple mental model

```
MCP Server
├─ Tools → Actions
├─ Resources → Data
└─ Prompts → Instructions
```

Q183. What is the difference between MCP Tool, Resource and Prompt?

MCP Component	Purpose
Tool	Perform an action
Resource	Provide/read information

MCP Component	Purpose
Prompt	Provide reusable instructions

EXAMPLE

For a GitHub system:

```
Tool:  
create_pull_request()  
Resource:  
repository://my-project  
Prompt:  
code_review
```

Q184. How does MCP work?

A simple flow:

```
User  
↓  
AI Agent  
↓  
LLM  
↓  
"I need GitHub information"  
↓  
MCP Client  
↓  
MCP Server  
↓  
GitHub  
↓  
Result  
↓  
MCP Client  
↓  
Agent  
↓  
LLM  
↓  
Final Answer
```

Important point

The LLM decides what capability is needed.

The MCP client communicates with the MCP server.

The MCP server interacts with the actual external system.

Q185. Does MCP replace APIs?

No.

This is an important interview answer.

MCP does not replace APIs. MCP provides a standardized way for AI applications to discover and use capabilities exposed by systems.

Think:

```
API
↓
Application-to-Application communication
MCP
↓
AI application ↔ tools/data
```

An MCP server may internally call APIs.

For example:

```
Agent
↓
MCP Client
↓
MCP Server
↓
REST API
↓
External System
```

1. MCP vs API — API

An API is an interface provided by a service.

EXAMPLE

```
GET /customers/1001
```

Your application knows how to call that API.

MCP

MCP standardizes how AI applications interact with capabilities.

```
AI Application
↓
MCP
↓
Tool
```

INTERVIEW ANSWER

"An API is a service interface. MCP is a protocol for exposing and consuming capabilities in a way that is designed for AI applications. An MCP server can itself use APIs internally."

1. MCP vs Function Calling — This is a very common interview question.

Function Calling

The model can request:

```
get_weather(city="Hyderabad")
```

Your application executes the function.

MCP

MCP provides a standardized protocol for discovering and interacting with tools and other server capabilities.

```
Agent
↓
MCP Client
↓
MCP Server
↓
Tool
```

SIMPLE ANSWER

Function calling is a model/application mechanism for requesting a function. MCP is a protocol for standardizing how AI applications connect to external tools and data.

They can work together.

1. MCP vs LangChain — These solve different problems.

LangChain

Framework for building LLM applications.

```
Prompt
LLM
Tools
RAG
Agents
Retrievers
```

MCP

Protocol for connecting AI applications to external capabilities.

```
AI Application
↓
MCP
↓
External Tools / Data
```

They can work together

```
LangChain Agent
  ↓
MCP Client
  ↓
MCP Server
  ↓
Database / API / GitHub
```

Experienced answer

"LangChain is an application framework, while MCP is an interoperability protocol. I can use LangChain to build the Agent and MCP to connect that Agent to standardized external tools."

1. MCP vs LangGraph — Again, they solve different problems.

LangGraph

Controls:

- State
- Nodes
- Edges
- Loops
- Routing
- Human approval
- Agent orchestration

MCP

Connects:

- Tools
- Resources
- Prompts
- External systems

So:

```
LangGraph
  ↓
Agent
  ↓
MCP Client
  ↓
MCP Server
  ↓
External System
```


INTERVIEW ANSWER

"LangGraph controls how the Agent workflow executes, while MCP controls how the Agent can connect to external capabilities."

Q186. Can an Agent use multiple MCP Servers?

Yes.

For example:




The Agent can use different capabilities from different MCP servers.

Q187. Can multiple Agents use one MCP Server?

Yes.

For example:



This is one reason standardization is useful.

Multiple AI applications or Agents can consume the same server capabilities, subject to authentication and authorization.

Q188. Can MCP connect to databases?

Yes.

For example:



The MCP server could expose tools such as:

```
get_customer()  
get_orders()  
search_products()
```

Important security point

Don't expose unrestricted SQL access unless there is a strong reason and proper controls.

Prefer:

```
Agent  
↓  
get_customer()  
↓  
Validation  
↓  
Database
```

rather than:

```
Agent  
↓  
"Run any SQL you want"  
↓  
Production Database
```

Q189. Can MCP connect to GitHub?

Yes.

Conceptually:

```
Agent  
↓  
MCP Client  
↓  
GitHub MCP Server  
↓  
GitHub
```

The server can expose operations such as:

```
search repositories  
read files  
create issues  
create pull requests
```

The exact available tools depend on the particular MCP server implementation.

Q190. Can MCP connect to AWS or Azure?

Yes.

You can create or use MCP servers that expose cloud capabilities.

For example:

```
Agent
↓
MCP Client
↓
AWS MCP Server
↓
AWS Services
```

or:

```
Agent
↓
MCP Client
↓
Azure MCP Server
↓
Azure Services
```

Important interview point

MCP does not magically provide access to AWS or Azure.

You need an MCP server implementation that exposes the required capabilities and appropriate credentials/permissions.

Q191. How do you build an MCP Server?

Conceptually, the steps are:

1. Choose SDK
- ↓
2. Create MCP Server
- ↓
3. Define Tools
- ↓
4. Define Resources
- ↓
5. Define Prompts if needed
- ↓
6. Add Authentication/Authorization
- ↓
7. Connect to External System
- ↓
8. Run/Deploy Server

Official MCP SDKs are available for languages including TypeScript, Python, Go and C#. (Model Context Protocol Blog)

Q192. Simple MCP Server Example

Conceptually, suppose we want a customer service MCP server.

```
Customer MCP Server
Tools:
  get_customer()
  get_order()
  cancel_order()
Resources:
  customer://profile
  customer://orders
Prompts:
  customer_support
```

Then:

```
AI Agent
  ↓
MCP Client
  ↓
Customer MCP Server
  ↓
Customer Database
```

The Agent can use the exposed capabilities.

Q193. How does MCP discovery work?

An MCP client needs to know what capabilities are available.

For example:

```
MCP Server
  ↓
Available Tools
  ├── search_customer
  ├── get_order
  └── cancel_order
```

The client can discover the available capabilities and their schemas.

The current 2026-07-28 specification also supports cacheable list results and an optional server/discover capability discovery RPC. (Model Context Protocol Blog)

Q194. What is the MCP transport?

Transport means:

How the client and server communicate.

MCP supports transport mechanisms such as:

```
stdio
HTTP-based transport
```

The current specification has moved toward a stateless HTTP-oriented protocol core; older HTTP+SSE approaches are deprecated in the 2026-07-28 specification. (Model Context Protocol Blog)

INTERVIEW ANSWER

"MCP defines the protocol and supports transports for communication. For local integrations, stdio is common; for remote production servers, HTTP-based transport is important."

Q195. What is stdio transport?

stdio means:

```
AI Application
  ↓
stdin/stdout
  ↓
MCP Server
```

It is commonly useful when the MCP server runs locally as a process.

EXAMPLE

```
VS Code
  ↓
MCP Client
  ↓
Local MCP Server
```

Q196. What is Remote MCP?

Instead of running the MCP server locally:

```
AI App
  ↓
Local MCP Server
```

we can have:

```
AI App
↓
Internet / Network
↓
Remote MCP Server
↓
Database / API / Cloud
```

This is important for enterprise deployment.

Q197. How do you secure MCP?

This is a very important experienced-level question.

I would consider:

Authentication

Who is connecting?

Authorization

What can they access?

Tool-level permissions

What can each tool do?

Input validation

Validate tool arguments.

Output validation

Validate important results.

Secrets

Don't expose API keys through prompts.

Least privilege

Give only required permissions.

Audit logging

Track important operations.

The current MCP specification has strengthened its authorization model, including issuer validation and credential binding. (Model Context Protocol Blog)

Q198. Why is security especially important with MCP?

Because MCP can expose real capabilities.

For example:

```
read_file()
delete_file()
send_email()
create_payment()
deploy_application()
```

If the Agent gets unrestricted access, it can potentially perform dangerous actions.

So:

```
Agent
↓
MCP Tool
↓
Authorization
↓
Validation
↓
Business Rules
↓
External System
```

Strong interview answer

"I treat MCP tools as privileged interfaces. I don't give an Agent unrestricted access just because the tool is available. I apply authentication, authorization, validation and least-privilege permissions."

Q199. Should every MCP tool be allowed to modify data?

No.

Separate tools based on risk.

Read tools

```
get_customer()
get_order()
search_product()
```

Usually lower risk.

Write tools

```
update_customer()
cancel_order()
send_email()
```

Need stronger controls.

High-risk tools

```
delete_database()
deploy_production()
transfer_money()
```

May require explicit approval.

Q200. What is Human-in-the-Loop with MCP?

Suppose an Agent wants to send an email.

```

Agent
↓
MCP send_email()
↓
Approval Required
↓
Human
↓
Approve
↓
Send Email


```

This is a good production architecture for sensitive actions.

Q201. Can MCP be used with LangGraph?

Yes.

A very common architecture can be:




LangGraph controls the workflow.

MCP provides external capabilities.

Q202. Can MCP be used with Multi-Agent systems?

Yes.

For example:



Different Agents can use different MCP servers.

This creates a clean separation between:

Agent responsibility and external capability.

Q203. MCP vs REST API — Real Example

Suppose we have a weather API.

Traditional API approach

Developer writes:

```

response = requests.get(
    "https://weather-api.com/weather?city=Hyderabad"
)
  
```

The application knows:

- URL
- HTTP method
- Authentication
- Parameters
- Response format

MCP approach

```

Agent
↓
MCP Client
↓
Weather MCP Server
↓
Weather API
  
```

The MCP server handles the underlying API details.

INTERVIEW ANSWER

"MCP can sit above existing APIs. The MCP server can translate a standardized tool request into the API call required by the underlying service."

Q204. Is MCP only for LLMs?

MCP is designed for AI applications and models interacting with external capabilities.

The important idea is not:

"MCP is an LLM."

It isn't.

MCP is a protocol.

```
LLM
↓
AI Application
↓
MCP
↓
Tools / Resources
```

Q205. Is MCP an Agent?

No.

This is very important.

```
Agent
↓
Uses MCP
↓
Tools
```

MCP does not make an application an Agent.

SIMPLE ANSWER

"MCP is a protocol. An Agent is an AI system that can reason and take actions. An Agent can use MCP to access external capabilities."

Q206. Is MCP a Vector Database?

No.

MCP is not:

- Vector DB
- LLM
- Agent
- RAG framework
- Database

It is a protocol.

Q207. What is the relationship between MCP and RAG?

They solve different problems.

RAG

Gets relevant information:

```
graph TD
    Question --> Retriever
    Retriever --> Documents
    Documents --> LLM
```

MCP

Connects the Agent to external capabilities:

```
graph TD
    Agent --> MCP
    MCP --> Tool_Resource[Tool / Resource]
```

They can be combined:


```
graph TD
    Agent --- RAG[--- RAG --> Company Documents]
    Agent --- MCP[--- MCP --> External Systems]
```

Q208. MCP + RAG Example

Suppose an HR Agent needs:

"Check the company leave policy and tell me how many leave days I have."

It could use:



This is a very realistic enterprise architecture.


Q209. Why is MCP important for Agentic AI?

Because Agents need access to the outside world.

Without external tools:



With MCP:



Experienced answer

"Agents become much more useful when they can interact with real systems. MCP provides a standardized way to expose those capabilities so AI applications can discover and use them."

Q210. What is the biggest advantage of MCP?

The key advantage is:

Standardization and interoperability.

Instead of building:

```
Agent → Custom GitHub Integration
Agent → Custom Slack Integration
Agent → Custom Database Integration
```

we can use:

```
Agent
↓
MCP
↓
Standardized Servers
```

This makes integrations more reusable.

Q211. What are the challenges of MCP?

Important challenges:

Security

Tools can perform real actions.

Authorization

Different users need different permissions.

Tool discovery

Too many tools can make selection difficult.

Reliability

Remote services can fail.

Latency

Network calls add latency.

Governance

Enterprises need to control which servers and tools are allowed.

INTERVIEW ANSWER

"MCP simplifies integration, but it doesn't remove the engineering challenges. In production I still need authentication, authorization, governance, validation, observability and failure handling."

Q212. What happens if an MCP Server goes down?

The Agent should not simply fail silently.

Possible architecture:

```
Agent
↓
MCP Server
↓
ERROR
↓
Retry?
↓
Fallback?
↓
Human?
↓
Final Error
```

For critical systems, use:

- Timeouts
- Retries
- Health checks
- Circuit breakers
- Fallbacks
- Monitoring

Q213. How would you deploy an MCP Server?

For production:

```
Code
↓
Docker
↓
Container Registry
↓
AWS / Azure
↓
Load Balancer
↓
MCP Server
↓
External Systems
```

AWS example

```
GitHub
↓
Jenkins
↓
Docker
↓
ECR
↓
ECS / EC2
↓
MCP Server
```

Azure example

```
GitHub
↓
Jenkins
↓
Docker
↓
Azure Container Registry
↓
Container Apps / App Service
↓
MCP Server
```

The exact choice depends on scale, networking, security and operational requirements.

Q214. Real Project Interview Question


Interviewer:

"How would you use MCP in your Agentic AI application?"

Strong answer:

"I would use MCP to standardize access to external capabilities. For example, if I build a customer-support Agent, I could expose customer lookup, order status and refund operations through an MCP server. The Agent would connect through an MCP client and select the required tool based on the user's request. I would keep authentication, authorization and validation at the tool and server boundaries. If the workflow is complex, LangGraph can orchestrate the Agent while MCP handles external tool integration."

ARCHITECTURE



Q215. "Why MCP instead of writing a normal tool?"

This is a good experienced-level question.

"If the tool is only needed inside one application, a normal function or internal API may be enough. I consider MCP when I want standardized, reusable access to tools or resources across multiple AI applications or Agents."


This is an important architectural distinction.

Don't say MCP is always better.

Q216. "Is MCP mandatory for Agentic AI?"

No.

You can build:



without MCP.

MCP becomes valuable when you need standardized and reusable integrations.

INTERVIEW ANSWER

"No. MCP is not mandatory for building Agents. It becomes useful when we want standardized interoperability between AI applications and external tools or resources."

Q217. Explain MCP in 30 Seconds

Memorize this:

"MCP stands for Model Context Protocol. It is an open standard that allows AI applications to connect with external tools, resources and services in a standardized way. An MCP server exposes capabilities, and an MCP client connects the AI application to that server. I can use MCP with LangChain or LangGraph. LangGraph handles Agent workflow and MCP handles external capability integration."


Q218. MCP — Quick Interview Revision

Question	Simple Answer
MCP?	Model Context Protocol
What is it?	Standard for connecting AI apps to tools/data
MCP Server?	Exposes tools, resources and prompts
MCP Client?	Connects AI app to MCP server
MCP Host?	AI application managing MCP connections
Tool?	Performs an action
Resource?	Provides/accesses information
Prompt?	Reusable instruction
MCP vs API?	MCP standardizes AI-facing integration; APIs expose services
MCP vs Function Calling?	Function calling requests functions; MCP standardizes external capability access
MCP vs LangChain?	Protocol vs application framework
MCP vs LangGraph?	External integration vs workflow orchestration
Can MCP use APIs?	Yes
Can MCP use databases?	Yes
Can MCP use GitHub?	Yes
Can multiple Agents use MCP?	Yes
Can one Agent use multiple MCP servers?	Yes
Is MCP an Agent?	No
Is MCP an LLM?	No
Is MCP a Vector DB?	No

Question	Simple Answer
Is MCP mandatory?	No
Main benefit?	Interoperability and standardization
Main challenge?	Security, governance, reliability and tool control

The Most Important MCP Architecture

Remember this diagram:



And remember the difference:

LangChain → Build LLM applications
LangGraph → Control Agent workflows
MCP → Connect Agents to external capabilities
LLM → Reason / generate
Tools → Perform actions
RAG → Retrieve knowledge

Quick Revision

Scan through these questions to self-test your recall before an interview.

- Q001. What is an AI Agent?
- Q002. What is the difference between an LLM and an AI Agent?
- Q003. What are the main components of an AI Agent?
- Q004. How does an AI Agent make decisions?
- Q005. What is the Agent Execution Loop?
- Q006. What is Tool Calling?
- Q007. What is Function Calling?
- Q008. What is the difference between Tool Calling and Function Calling?
- Q009. What is Memory in an AI Agent?
- Q010. What is Short-Term Memory?
- Q011. What is Long-Term Memory?
- Q012. What is Planning in Agentic AI?
- Q013. What is Reasoning in an AI Agent?
- Q014. What is Reflection?
- Q015. What is ReAct?
- Q016. What is the difference between a Workflow and an Agent?
- Q017. When should we use an Agent instead of a normal LLM application?
- Q018. What are the limitations of AI Agents?
- Q019. How do you prevent an Agent from entering an infinite loop?
- Q020. How do you handle hallucinations in Agents?
- Q021. How do you evaluate an AI Agent?
- Q022. What is an Agent Loop?
- Q023. What is State in an Agent?
- Q024. What is Agent Orchestration?
- Q025. How do you design a production-ready AI Agent?
- Q026. LangChain — Interview Questions
- Q027. What is LangChain?
- Q028. Why do we need LangChain?
- Q029. What are the main components of LangChain?
- Q030. What is a Prompt Template?
- Q031. What is an LLM in LangChain?
- Q032. What is a Chain?
- Q033. What is LCEL?
- Q034. What is Runnable?
- Q035. What is a Tool in LangChain?
- Q036. What is an Agent in LangChain?
- Q037. What is the difference between Chain and Agent?
- Q038. How does a LangChain Agent work?
- Q039. What is a Retriever?
- Q040. What is RAG in LangChain?
- Q041. What is a Document Loader?
- Q042. What is Text Splitting?
- Q043. Why?

- Q044. What is an Embedding?
- Q045. What is a Vector Store?
- Q046. What is Memory in LangChain?
- Q047. How do you connect a database with LangChain?
- Q048. How do you connect an API with LangChain?
- Q049. What is Structured Output?
- Q050. How do you handle errors in LangChain?
- Q051. LangChain vs LangGraph?
- Q052. Is LangChain mandatory for building an AI Agent?
- Q053. When would you NOT use LangChain?
- Q054. How would you make a LangChain application production-ready?
- Q055. How would you explain LangChain in an interview in 30 seconds?
- Q056. LangGraph — Interview Questions & Answers
- Q057. What is LangGraph?
- Q058. Why do we need LangGraph?
- Q059. What is a Graph in LangGraph?
- Q060. What is a Node?
- Q061. What is an Edge?
- Q062. What is a Conditional Edge?
- Q063. What is State in LangGraph?
- Q064. Why is State important?
- Q065. What are START and END?
- Q066. How does a LangGraph workflow execute?
- Q067. What is the difference between LangChain and LangGraph?
- Q068. Can LangGraph work without LangChain?
- Q069. What is a LangGraph Agent?
- Q070. What is a loop in LangGraph?
- Q071. Why are loops useful in Agentic AI?
- Q072. What is Human-in-the-Loop?
- Q073. What is Checkpointing?
- Q074. Why is persistence important?
- Q075. How do you handle errors in LangGraph?
- Q076. How do you prevent infinite loops?
- Q077. How do you build a Multi-Agent system using LangGraph?
- Q078. What is a Supervisor Agent?
- Q079. What is a Subgraph?
- Q080. Why is LangGraph useful for production?
- Q081. How would you design a Customer Support Agent using LangGraph?
- Q082. How would you design an Agentic RAG system using LangGraph?
- Q083. LangChain Agent vs LangGraph Agent?
- Q084. How do you debug a LangGraph application?
- Q085. How do you control Agent behavior in LangGraph?
- Q086. What are the advantages of LangGraph?
- Q087. What are the disadvantages/challenges?
- Q088. When should you NOT use LangGraph?
- Q089. Explain LangGraph in 30 Seconds
- Q090. Explain LangGraph using a Real Project

- Q091. Single Agent — Interview Questions & Answers
- Q092. What is a Single Agent?
- Q093. How is a Single Agent different from a normal chatbot?
- Q094. What are the main components of a Single Agent?
- Q095. How does a Single Agent work?
- Q096. How does the Agent decide which tool to use?
- Q097. Can a Single Agent use multiple tools?
- Q098. Can a Single Agent use multiple tools in one task?
- Q099. What is the Agent Loop?
- Q100. Does the Agent execute tools itself?
- Q101. What is the role of the LLM in a Single Agent?
- Q102. What is the role of tools?
- Q103. What is the role of memory?
- Q104. What is State in a Single Agent?
- Q105. When should you use a Single Agent?
- Q106. Why not use Multi-Agent for everything?
- Q107. Single Agent vs Multi-Agent
- Q108. What are the advantages of a Single Agent?
- Q109. What are the limitations of a Single Agent?
- Q110. What is Tool Selection?
- Q111. How do you prevent an Agent from using the wrong tool?
- Q112. How do you handle a dangerous tool?
- Q113. How do you prevent infinite loops in a Single Agent?
- Q114. How do you handle tool failures?
- Q115. How do you handle hallucinations in a Single Agent?
- Q116. How do you evaluate a Single Agent?
- Q117. How many tool calls?
- Q118. How do you monitor a Single Agent in production?
- Q119. How do you secure a Single Agent?
- Q120. Can a Single Agent use RAG?
- Q121. Can a Single Agent use MCP?
- Q122. Can a Single Agent be built using LangGraph?
- Q123. Real Project Interview Question
- Q124. "Why did you choose Single Agent?"
- Q125. "When would you move from Single Agent to Multi-Agent?"
- Q126. Explain Single Agent in 30 Seconds
- Q127. Single Agent — Quick Revision
- Q128. Multi-Agent Systems — Interview Questions & Answers
- Q129. What is a Multi-Agent System?
- Q130. Why do we need Multi-Agent Systems?
- Q131. What is an Agent's role in a Multi-Agent system?
- Q132. What is Agent Specialization?
- Q133. What is a Supervisor Agent?
- Q134. What is an Orchestrator?
- Q135. How do Agents communicate with each other?
- Q136. What is Shared State?
- Q137. What is Agent Delegation?

- Q138. What are the common Multi-Agent architectures?
- Q139. What is the Supervisor Pattern?
- Q140. What is Sequential Multi-Agent Architecture?
- Q141. What is Parallel Multi-Agent Architecture?
- Q142. What is Hierarchical Multi-Agent Architecture?
- Q143. What is Peer-to-Peer Agent Architecture?
- Q144. Single Agent vs Multi-Agent
- Q145. Why not simply give all tools to one Agent?
- Q146. How do you decide whether to use Single Agent or Multi-Agent?
- Q147. What are the advantages of Multi-Agent Systems?
- Q148. What are the disadvantages?
- Q149. How do you prevent Agents from duplicating work?
- Q150. How do you prevent Agent conflicts?
- Q151. Who decides?
- Q152. Should Agents directly modify databases?
- Q153. How do you secure a Multi-Agent system?
- Q154. How do Agents share information?
- Q155. How would you implement Multi-Agent using LangGraph?
- Q156. Can Multi-Agent systems use MCP?
- Q157. Multi-Agent + MCP Example
- Q158. How do you handle failures in Multi-Agent systems?
- Q159. How do you evaluate a Multi-Agent system?
- Q160. What is Agent-to-Agent communication?
- Q161. What is Agent Handoff?
- Q162. What is a Shared-State Multi-Agent system?
- Q163. What is a Message-Based Multi-Agent system?
- Q164. What is the biggest challenge in Multi-Agent systems?
- Q165. Real Project Interview Question
- Q166. How would you answer: "Why did you choose Multi-Agent?"
- Q167. How would you answer: "Why didn't you use Multi-Agent?"
- Q168. How do you reduce cost in Multi-Agent systems?
- Q169. How do you reduce latency?
- Q170. Single Agent vs Multi-Agent — Interview Table
- Q171. Explain Multi-Agent in 30 Seconds
- Q172. MCP — Model Context Protocol
- Q173. What is MCP?
- Q174. Why was MCP introduced?
- Q175. What problem does MCP solve?
- Q176. What is an MCP Server?
- Q177. What is an MCP Client?
- Q178. What is an MCP Host?
- Q179. MCP Host vs Client vs Server
- Q180. What is an MCP Tool?
- Q181. What is an MCP Resource?
- Q182. What is an MCP Prompt?
- Q183. What is the difference between MCP Tool, Resource and Prompt?
- Q184. How does MCP work?

- Q185. Does MCP replace APIs?
- Q186. Can an Agent use multiple MCP Servers?
- Q187. Can multiple Agents use one MCP Server?
- Q188. Can MCP connect to databases?
- Q189. Can MCP connect to GitHub?
- Q190. Can MCP connect to AWS or Azure?
- Q191. How do you build an MCP Server?
- Q192. Simple MCP Server Example
- Q193. How does MCP discovery work?
- Q194. What is the MCP transport?
- Q195. What is stdio transport?
- Q196. What is Remote MCP?
- Q197. How do you secure MCP?
- Q198. Why is security especially important with MCP?
- Q199. Should every MCP tool be allowed to modify data?
- Q200. What is Human-in-the-Loop with MCP?
- Q201. Can MCP be used with LangGraph?
- Q202. Can MCP be used with Multi-Agent systems?
- Q203. MCP vs REST API — Real Example
- Q204. Is MCP only for LLMs?
- Q205. Is MCP an Agent?
- Q206. Is MCP a Vector Database?
- Q207. What is the relationship between MCP and RAG?
- Q208. MCP + RAG Example
- Q209. Why is MCP important for Agentic AI?
- Q210. What is the biggest advantage of MCP?
- Q211. What are the challenges of MCP?
- Q212. What happens if an MCP Server goes down?
- Q213. How would you deploy an MCP Server?
- Q214. Real Project Interview Question
- Q215. "Why MCP instead of writing a normal tool?"
- Q216. "Is MCP mandatory for Agentic AI?"
- Q217. Explain MCP in 30 Seconds
- Q218. MCP — Quick Interview Revision